

CARSKit: A Java-Based Context-aware Recommendation Engine

Yong Zheng

Center for Web Intelligence
DePaul University
Chicago, Illinois, 60604, USA
yzheng8@cs.depaul.edu

Bamshad Mobasher

Center for Web Intelligence
DePaul University
Chicago, Illinois, 60604, USA
mobasher@cs.depaul.edu

Robin Burke

Center for Web Intelligence
DePaul University
Chicago, Illinois, 60604, USA
rburke@cs.depaul.edu

Abstract—Recommender system has been demonstrated as one of the most useful tools to assist users’ decision makings. Several recommendation algorithms have been developed and implemented by both commercial and open-source recommendation libraries. Context-aware recommender system (CARS) emerged as a novel research direction during the past decade and many contextual recommendation algorithms have been proposed. Unfortunately, no recommendation engines start to embed those algorithms in their kits, due to the special characteristics of the data format and processing methods in the domain of CARS. This paper introduces an open-source Java-based context-aware recommendation engine named as CARSKit which is recognized as the 1st open source recommendation library specifically designed for CARS. It implements the state-of-the-art context-aware recommendation algorithms, and we will showcase the ease with which CARSKit allows recommenders to be configured and evaluated in this demo.

I. INTRODUCTION AND MOTIVATIONS

Recommender systems (RS) have been used to successfully address the information overload problem by providing recommendations to the end users. In recent years, researchers have realized the importance of contexts and try to incorporate contexts into RS to build context-aware recommender systems (CARS) which are able to adapt recommendations to users’ contextual situations.

Context is defined as any information that can be used to characterize any entity [4]. In CARS, contexts are usually selected from the dynamic attributes of the activity itself [9]. For example, user may choose a different type of movie if he or she is going to watch the movie *with kids* rather than *with partners*, where *companion* is one attribute from the activity “watching a movie”, and its value may change frequently when the same activity happens repeatedly. In addition, limited attributes from user profiles can be considered as contexts too, such as user intent and emotional states. Most item features, such as movie genre, music singer, or hotel room services, are considered as content in RS.

CARS have been well developed in the past decade. Many classical contextual recommendation algorithms, such as context-aware splitting approaches (CASA) [3], [10], tensor factorization (TF) [6] and context-aware matrix factorization (CAMF) [2], have been well recognized as the standard baselines in this domain. However, there are no recommendation engines implementing context-aware recommendation algorithms, due to the special characteristics of the data format

and processing methods in this domain. Researchers have to compile those algorithms by themselves again and again, which further impedes the development progress of CARS.

Therefore, it is badly in need of such a recommendation library to not only meet the basic design and evaluation purpose, but also boost the further development of context-aware recommendations. We introduce CARSKit – an open-source Java-based context-aware recommendation engine which implements the state-of-the-art contextual recommendation algorithms and provides a standard platform for evaluations and practical use. As far as we know, it is the 1st open source library specifically designed for CARS and it is able to provide flexible configurations and compatible spaces for expansion with new algorithms.

II. RELATED WORK

There are at least three reasons to develop a recommendation engine: 1). classical recommendation algorithms are the frequent ones to be evaluated in practice. A recommendation library is able to provide standard implementations and avoid repeatedly compiling those algorithms from time to time; 2). it is also a standard platform for benchmark or evaluations; 3). it is also helpful for both research purpose and industrial practice, as well as educational tools in teaching and learning.

During the past decades, tons of recommendation algorithms have been proposed and several recommendation engines, such as Mahout¹, Duine², Cofi³, EasyRec⁴, GraphLab Create⁵, LensKit⁶, LibRec⁷, MyMediaLite⁸, started to embed those algorithms into their kits. Here, we select the top popular ones and provide a brief comparison shown in Table I.

Mahout and GraphLab Create are two large-scale data mining and machine learning platforms, where GraphLab was an open-source library, but currently it was developed as one of DATO⁹ commercial products. MyMediaLite was one of the most popular recommendation libraries but it was no

¹ Mahout, <http://mahout.apache.org/>

² Duine, <http://www.duineframework.org/>

³ Cofi, <http://www.nongnu.org/cofi/>

⁴ EasyRec, <http://easyrec.org/>

⁵ GraphLab Create, <https://dato.com/products/create/>

⁶ LensKit, <http://lenskit.org/>

⁷ LibRec, <http://www.librec.net/>

⁸ MyMediaLite, <http://www.mymedialite.net/>

⁹ DATO, <https://dato.com/>

TABLE I: Comparison Among Selected Recommendation Engines

	Mahout	GraphLab Create	MyMediaLite	EasyRec	LensKit	LibRec	CARSKit
Initial release (date)	04/2007	2014	10/2010	03/2011	03/2011	2014	09/2015
Latest release (date)	05/2015	2015	09/2013	05/2013	05/2015	06/2015	09/2015
Latest version	0.10.1	1.4.1	3.10	0.98	2.2	1.3.2	0.1.0
Active updates	Yes	Yes	No	No	Yes	Yes	Yes
License	Apache	BSD, AGPL	GPL	N/A	GPL	GPL	GPL
Programming language	Java	Python, C++	C#	WebService	Java	Java	Java
Running platform	JVM	Any	.NET	WebApp	JVM	JVM	JVM
Command line	Yes	Yes	Yes	No	Yes	Yes	Yes
Documentation	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Online updates	Yes	Yes	Yes	No	No	No	No
Algorithm diversity	Classical	Classical	State-of-the-art	Classical	Classical	State-of-the-art	CARS
Distributed computing	Partial	Yes	No	No	No	No	No
Open source	Yes	Partial	Yes	Yes	Yes	Yes	Yes

longer updated for a while, where Rival¹⁰ and WrapRec¹¹ were recently designed as a wrapper of MyMediaLite for standard benchmark and evaluations. They aim to provide a simple and unified way to create evaluation results comparable across different recommendation frameworks. EasyRec is a special one designed as a Web service deliver to the Web applications, but it was outdated and not updated anymore. LensKit and LibRec are two modern java-based recommendation libraries, where LensKit only provides the implementation of classical recommendation algorithms, such as user-based and item-based k NN collaborative filtering, slope one recommender, and matrix factorization. In fact, only MyMediaLite and LibRec are able to provide the state-of-the-art recommendation algorithms in addition to those classical ones, and LibRec runs faster based on a benchmark evaluation [5].

In contrast to those existing recommendation engines, CARSKit is specifically designed for context-aware recommendations, where the state-of-the-art recommendation algorithms in CARS are implemented in a way of easy-configuration and flexible-expansion.

III. CARSKIT: DESIGN AND FEATURES

CARSKit¹² is an open-source free software, where it can be used, modified and distributed under the terms of the GNU General Public License. In this section, we will introduce the specific design and features in CARSKit.

A. Architecture and Design

CARSKit provides a flexible architecture so that it is easy to expand the scope of context-aware recommendation algorithms and provides spaces to develop new algorithms in the future. The whole design can be depicted by the Figure 1.

The workflow is straightforward in our design: different recommendation algorithms are the specific implementations and extensions from the generic interfaces where the shared and common functions are defined, such as rating or score prediction for a user on one item in a specific context. Evaluations for rating predictions and top- N recommendations are embedded into the *Recommender*.

B. Features

Due to the special characteristics of the data format and processing methods in CARS, there are several challenges in the design. In this section, we will introduce the key features specifically designed for CARS, which can further be stated from those aspects: data transformer, data structure, recommendation algorithms, configuration and evaluations.

1) *Data Transformer*: In traditional RS, the rating prediction problem begins with a two dimensional matrix of ratings: $Users \times Items \rightarrow Ratings$. In CARS, contexts are considered as additional inputs, which results in a multidimensional rating function: $Users \times Items \times Contexts \rightarrow Ratings$ [1]. This special data format brings challenges to build such a recommendation kit, where the original structure in other recommendation engines may not be easily reused for context-aware recommendation.

Before the discussion of data structure, we introduce the data transformer here. Usually, the contextual rating data can be stored in two formats: *loose* format and *compact* format, as shown in tables below.

UserID	ItemID	Rating	Context	Condition
U1	T1	3	Time	Weekend
U1	T1	3	Location	Work
U2	T2	4	Time	Weekday
U2	T2	4	Location	Home

TABLE II: Loose Format

UserID	ItemID	Rating	Time	Location
U1	T1	3	Weekend	Work
U2	T2	4	Weekday	Home
U1	T1	4	Weekend	Home
U2	T2	2	Weekday	Work

TABLE III: Compact Format

Context dimension is identical to contextual variable, e.g., *Time* and *Location* as shown in the example above. Context conditions refer to specific values in a dimension, e.g. *Weekend* and *Weekday* are two conditions in the dimension *Time*. The loose format assumes that there is only one rating for each $\langle \text{user}, \text{item} \rangle$ pair in associated contexts, where the compact format allows to store multiple ratings to a same $\langle \text{user}, \text{item} \rangle$ pair in different contextual situations. Take the example shown in the two tables above, the first two rows in loose format actually represent a single rating by U1 for T1 within contexts {Weekend, Work}. In compact format, each row represents a single contextual rating profile; that is, there are only two

¹⁰Rival, <http://rival.recommenders.net/>

¹¹WrapRec, <https://github.com/babakx/WrapRec/>

¹²CARSKit, <https://github.com/ircsys/CARSKit/>

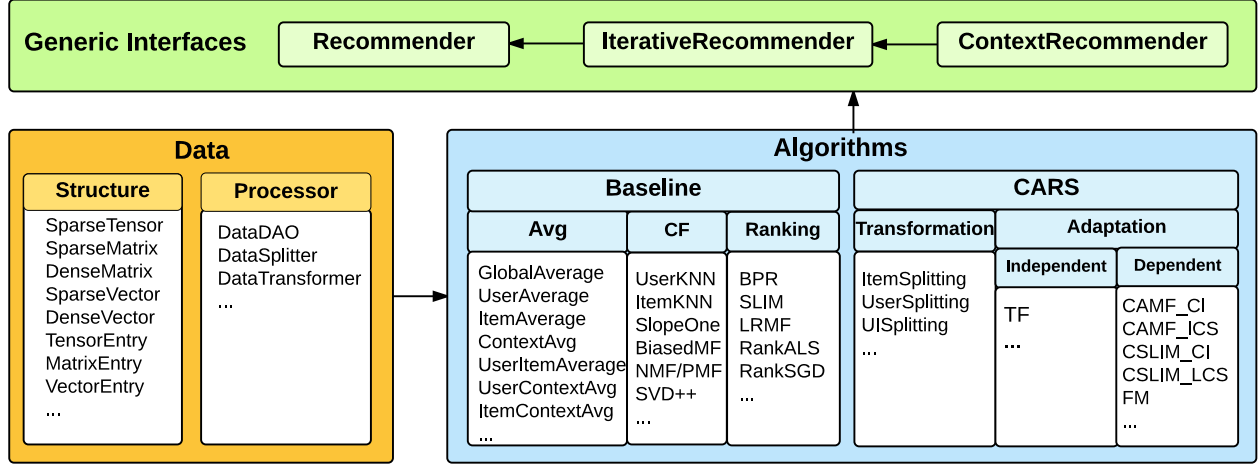


Fig. 1: CARSKit: The Architecture and Design

contextual rating profiles in the loose format but four rating profiles in the compact format in this example.

UserID	ItemID	Rating	Time:Weekend	Time:Weekday	Location:Home	Location:Work
U1	T1	3	1	0	0	1
U2	T2	4	0	1	1	0
U1	T1	4	1	0	1	0
U2	T2	2	0	1	0	1

TABLE IV: Binary Format

Most contextual information is in shape of categorical data. Both the loose and compact format will increase storage pressure and computational costs. In CARSKit, we store contextual rating in a *binary* format as shown in Table IV, which is able to significantly boost the running performance. To assist the end users to prepare the rating data, we provide two methods *TransformationFromLooseToBinary* and *TransformationFromCompactToBinary* as the data transformer in our toolkit.

2) *Data Structure*: Since contexts are considered as additional inputs, the data structure becomes the most important part which directly leaves an impact on the flexibility and running performance.

There are two factors involved in designing the data structure: data storage and data operations. Intuitively, the context-aware data can be represented in a N -dimensional space as tensors, where each context dimension is considered as an individual dimension in the rating space. In this case, we build *SparseTensor* and *TensorEntry* to record the indices of each user, item, context dimension, and the associated rating. This structure is good for the context-aware recommendation algorithms using N -dimensional operations, such as the multiverse recommendation algorithm described in [6].

Meantime, there are many more contextual recommendation algorithms exploring the dependencies between contexts and user/item dimensions, where two-dimensional operation is still the most frequent one adopted in those algorithms. In this case, we continue to use *SparseMatrix* and *SparseVector* which are the well-recognized and good-efficient data representations in existing recommendation libraries, such as LensKit and LibRec. *SparseMatrix* uses the compressed row and column

storage¹³ which was also demonstrated to boost the running efficiency in the design of LibRec [5].

3) *Recommendation Algorithms*: As mentioned before, CARSKit is the 1st library specifically designed for CARS. As shown in Figure 1, we divide the contextual algorithms into two categories: transformation algorithms and adaptation algorithms.

The transformation algorithms try to pre-process the data and convert the contextual data set to a 2-dimensional rating matrix which only contains users, items and ratings, so that any traditional recommendation algorithms can be applied to. One of the most effective techniques falling into this category is the context-aware splitting approaches [3], [10].

The adaptation algorithms directly incorporate contexts into the prediction function. There are two subcategories involved: *independent modeling* (e.g., TF [6]) which assumes contexts are independent with users (and items), and *dependent modeling* which exploits the dependencies among users, items and contexts, such as CAMF [2] and contextual sparse linear method (CSLIM) [12], [13]. Dependent modeling can be built in two ways: by modeling contextual rating deviations [2], [13] and by learning context similarities [14], [15]. Factorization machines (FM) [7] is a finer-grained algorithm which exploits pairwise relationships in its learning process. Among those algorithms, TF and CAMF are two popular ones which have been recognized as the standard baselines in CARS.

In addition to those state-of-the-art contextual recommendation algorithms, we also include some traditional recommendation algorithms in the package *baseline*. We did not re-compile those algorithms and directly reuse the classical recommenders provided by LibRec. There are two main purposes to include those traditional recommendation algorithms – On one hand, those algorithms can be applied after the data transformation (e.g., splitting operations), which is an essential step in the context-aware transformation algorithms. On the other hand, it is usually common to compete a contextual recommendation algorithm with non-contextual algorithms to

¹³Sparse Matrix Storage, http://netlib.org/linalg/html_templates/node90.html

judge whether the contextual effect is significant or a context-aware recommendation algorithm is necessary or not.

4) *Configuration and Evaluations*: For evaluation purpose, we provide *DataSplitter* which enables the users to adopt either train-testing evaluation or the N -folds cross validations.

Most of the recommendation algorithms embedded in CARSKit are able to perform the two recommendation task: *rating prediction* and *item recommendation*, except those ones specifically designed for top- N recommendation, such as C-SLIM. But the evaluation is different from traditional ones, since contexts are additional inputs in the evaluation process. Typically, the rating prediction can be evaluated by different prediction errors, such as mean absolute error (MAE), root mean square error (RMSE) and mean prediction error (MPE). The item recommendation can be evaluated through *relevance* metrics, such as precision and recall, and *ranking* metrics, such as mean average precision (MAP), normalized discounted cumulative gain (NDCG) and mean reciprocal rank (MMR).

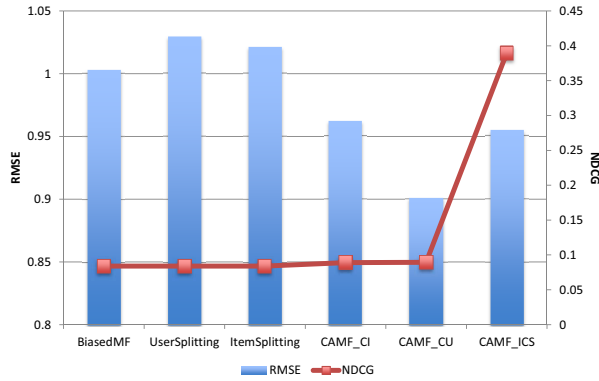


Fig. 2: Sample of Evaluations on *DePaulMovie* Data

We simply explored a 5-fold cross validation on *DePaulMovie* data¹⁴ by using recommendation algorithms built upon matrix factorization. The evaluation results based on rating predictions (by RMSE) and top-5 recommendations (by NDCG) are shown in Figure 2. Among of those algorithms, BiasedMF is a standard matrix factorization technique and also a non-contextual algorithm applied on the context-aware data. UserSplitting and ItemSplitting are two contextual splitting approaches [10] where BiasedMF is applied after data transformation. In addition, deviation-based CAMF (e.g., CAMF_CI and CAMF_CU) [2] and similarity-based CAMF (e.g., CAMF_ICS) [15] are also considered. The results reveal that CAMF_ICS is able to help obtain the best NDCG on this data set, while there are no statistical significant differences on RMSE among those algorithms.

Moreover, CARSKit provides flexible configurations by a single file which includes both *algorithm configuration* (e.g., algorithm parameters) and *experimental configuration*, e.g., inputs and output, evaluation strategy and metrics, etc.

IV. CONCLUSIONS AND FUTURE WORK

In this paper, we introduced CARSKit which is an open-source Java-based context-aware recommendation engine. CARSKit

contributes to the domain of recommender systems by providing 1). the 1st library specifically designed for context-aware recommendation; 2). implementations of the state-of-the-art contextual recommendation algorithms; 3). efficient data structure and flexible configuration, as well as multi-aspect evaluations. In future, we will add more state-of-the-art context-aware recommendation algorithms to CARSKit. It is also possible to include context suggestion algorithms, where context suggestion or context recommendation [11], [8] is a new direction derived from CARS which aims to recommend appropriate contexts for users to consume the items.

V. ACKNOWLEDGEMENT

We would like to show our gratitude to Dr. Guibing Guo (the author of LibRec) for his comments and suggestions on the development of CARSKit.

REFERENCES

- [1] G. Adomavicius, B. Mobasher, F. Ricci, and A. Tuzhilin. Context-aware recommender systems. *AI Magazine*, 32(3):67–80, 2011.
- [2] L. Baltrunas, B. Ludwig, and F. Ricci. Matrix factorization techniques for context aware recommendation. In *Proceedings of the fifth ACM conference on Recommender systems*, pages 301–304. ACM, 2011.
- [3] L. Baltrunas and F. Ricci. Experimental evaluation of context-dependent collaborative filtering using item splitting. *User Modeling and User-Adapted Interaction*, pages 1–28, 2013.
- [4] A. K. Dey. Understanding and using context. *Personal and ubiquitous computing*, 5(1):4–7, 2001.
- [5] G. Guo, J. Zhang, Z. Sun, and N. Yorke-Smith. Librec: A java library for recommender systems. In *Posters, Demos, Late-breaking Results and Workshop Proceedings of the 23rd International Conference on User Modeling, Adaptation and Personalization*, 2015.
- [6] A. Karatzoglou, X. Amatriain, L. Baltrunas, and N. Oliver. Multiverse recommendation: n-dimensional tensor factorization for context-aware collaborative filtering. In *Proceedings of the fourth ACM conference on Recommender systems*, pages 79–86. ACM, 2010.
- [7] S. Rendle, Z. Gantner, C. Freudenthaler, and L. Schmidt-Thieme. Fast context-aware recommendations with factorization machines. In *Proceedings of the 34th international ACM SIGIR conference on Research and development in Information Retrieval*, pages 635–644. ACM, 2011.
- [8] Y. Zheng. Context suggestion: Solutions and challenges. In *Proceedings of the 15th IEEE International Conference on Data Mining Workshops*, 2015.
- [9] Y. Zheng. A revisit to the identification of contexts in recommender systems. In *Proceedings of the 20th ACM Conference on Intelligent User Interfaces Companion*, pages 133–136. ACM, 2015.
- [10] Y. Zheng, R. Burke, and B. Mobasher. Splitting approaches for context-aware recommendation: An empirical study. In *Proceedings of the 29th ACM Symposium on Applied Computing*, pages 274–279. ACM, 2014.
- [11] Y. Zheng, B. Mobasher, and R. Burke. Context recommendation using multi-label classification. In *Proceedings of the 13th IEEE/WIC/ACM International Conference on Web Intelligence*, pages 288–295, 2014.
- [12] Y. Zheng, B. Mobasher, and R. Burke. CSLIM: Contextual SLIM recommendation algorithms. In *Proceedings of the 8th ACM Conference on Recommender Systems*, pages 301–304. ACM, 2014.
- [13] Y. Zheng, B. Mobasher, and R. Burke. Deviation-based contextual SLIM recommenders. In *Proceedings of the 23rd ACM Conference on Information and Knowledge Management*, pages 271–280, 2014.
- [14] Y. Zheng, B. Mobasher, and R. Burke. Integrating context similarity with sparse linear recommendation model. In *User Modeling, Adaptation, and Personalization*, pages 370–376. Springer, 2015.
- [15] Y. Zheng, B. Mobasher, and R. Burke. Similarity-based context-aware recommendation. In *Web Information Systems Engineering*. Springer, 2015.

¹⁴There are 5,029 ratings (scale 1-5) by 97 users on 79 movies within contexts “time, location, companion”.