



Product recommendations enhanced with reviews



Muthusamy Chelliah

Director, Academic Engagement, Flipkart

Sudeshna Sarkar

Professor, CSE, IIT Kharagpur

Email: sudeshna@cse.iitkgp.ernet.in



Reviews

Kamaran 8 May 2017

Denmark

9 reviews

9.2 "Best views of Como for a romantic stay"

- Breakfast could have been better, limited variety of things.
- + Impressive views of Como lake. Comfortable room. Friendly staff. Just 5 min walking to the boat station.

Helpful 1 person found this review helpful.

Which aspect is more important?

Food or Location?

What value of the aspect is desired?

Location near the lake or

Location near the conference venue?



ADD TO CART

BUY NOW

k to top

15,900 + ₹798 = ₹16,698

CART

4.1 ★



42,844

Ratings

10,516

Reviews

[Rate and Review product](#)
[Most Helpful](#) [Positive](#) [Negative](#) [Most Recent](#) [By Certified Buyers](#)

4★ Decent Performance

Pros:

- Decent camera - good in day light
- Battery lasts a business day with full 4G on
- Fingerprint scanner works well

Cons:

- Occasionally hangs
- Bit expensive for these feature set. Worth buy at 15k.

Ram N Certified Buyer 19 Dec, 2016

1729

200

2★ Expected a better product

Specs n performance r ok ...but d phone is too fragile ...a single drop from a vry height caused its display to shatter ...d display cost me a lil under 8k to replace at

Product Recommendation

The screenshot shows a product recommendation section on the Flipkart website. At the top, there's a search bar with the placeholder "Search for Products, Brands and More". Below it, a "Similar products" section displays two cameras: a Canon EOS 1300D and a Nikon D5200. Each camera has its name, a small image, a rating (4.4★ and 4.5★ respectively), the number of reviews (8,235 and 4,399), the original price (₹29,995 and ₹27,999), and the discounted price (₹22,495 and ₹27,999). A "25% off" badge is also visible. Below this, a "You may also be interested in" section shows two camera lenses.

- Suggest new products for current selection
 - Substitute
 - Complementary
- Help offer
 - Up-sell
 - Cross-sell
 - Bundle

Rich content in user reviews

About Looks : Very very impressive like a premium phone but back side is not so attractive, it's look a like Redmi Series Phone's as well as back camera edge style is boring.

I'm not impress with the picture quality coz its not what Samsung provide even in his low range phones, if you compare with moto series phones i will give 10/7 to samsung and moto 10/9, also the shake reduction rate on this phone is very bad, After taking the picture (naturally our hand shake little bit) the picture is blurry or shaked.

Sound : Not for who listening the music on speakers, While ringing it's good but not excellent

Well I'm shocked in packaging too as soon as i see the samsung box it's look they given me a 2 year old box

User Modeling

Recommendation Generation



Review Enhanced Recommendation

1. Introduction
2. Review mining
 - a) Background: Features and Sentiment, Latent features
 - b) Statistical approaches. LDA
 - c) Recent topic models
 - d) Deep Learning (ABSA)
 - e) Flipkart ABSA/Recommendation
3. Review based Recommendation
 - a) Handling Data Sparsity
 - b) Topic Model, Matrix Factorization and Mixture Model based Recommendation
 - c) Deep learning based Recommendation
4. Explainable Product Recommendations
5. Summary, Future Trend



Recommender System

Reviews

INTRODUCTION TO RECOMMENDER SYSTEMS



The Recommendation Task

- Rating Prediction: Estimate a utility function to predict how a user will like an item
- Recommendation: Recommend a set of items to maximize user's utility



Recommender Systems

Data

- Content
 - Product Information
 - Product Taxonomy
 - Product Attributes
 - Product Description
 - Customer Demographics
- User Activity
 - Purchase
 - Rating, Preference
 - Click / Browse
 - User generated reviews
- Social data

Methods

- Content based
 - Recommend items similar to those a user has liked in the past
 - Issues: Does not use quality judgments of users
- Collaborative filtering
 - Finds users with similar tastes as target user to make recommendation
 - Use ratings
 - Issues: Sparsity, Cold Start

Content based recommender systems

- Each item defined by a profile vector obtained from the item's textual description, metadata, etc.
- Weighting methods such as TF-IDF
- User profile vector : by aggregating the profile vectors of the items that the user liked or purchased in the past.
- Recommend items with high similarity between user vector and item vector.

Product description



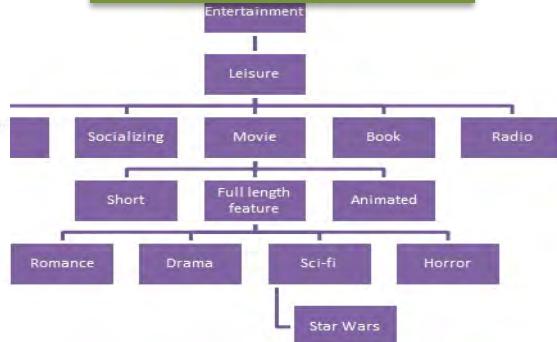
Lenovo A850

1 GB RAM

4.1 ★ 138 Ratings & 43 Reviews

- 1 GB RAM | 4 GB ROM | Expandable Upto 32 GB
- 5.5 inch quarter HD Display
- 5MP Rear Camera | 0.3MP Front Camera
- 2250 mAh Li-Polymer Battery
- MTK 6582M Processor
- 1 Year Manufacturer Warranty

Product Ontology



User Review

4 ★ Good choice

Excellent product.Fast delivery by Ekart.There was no colour selection,only Alfa steel was available in flipkart.Installation through Flipkart takes minimum two days, but on contacting the local whirlpool centre they installed it on the same day of delivery.Thanks to whirlpool flipkart and Ekart.The product is very useful only drawback is no lighting in the top deepfreezer compartment.Anyway satisfied with this purchase.

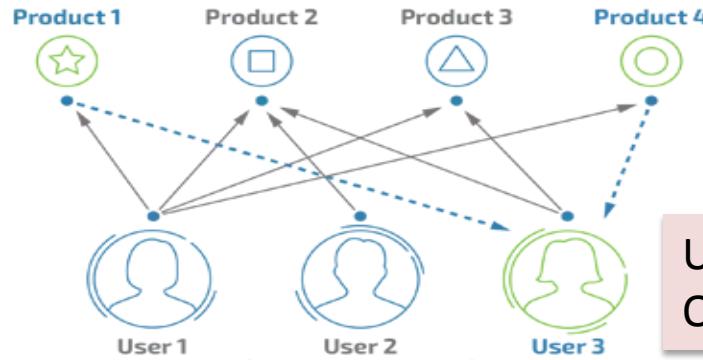
Damodaran Janardanan Certified Buyer 27 Dec, 2016

137 9

Collaborative filtering

Predict user preferences on new items by collecting rating information from many users (collaborative).

- Many items to choose from
- Few data per user
- No data for new user



Difficulties

- Data sparsity
 - too few ratings
 - 90%+ items with less than 10 reviews
- Cold-start
 - New user/item appear continuously

User Reviews provide additional rich data.
One Review worth much more than one rating.

Collaborative filtering

Memory Based CF:

1. User-based CF
2. Item-based CF

Model-based CF:

- Train a parametric model with the rating matrix
- Applied to predict ratings of unknown items
 - Ex: latent factor model



Latent Factor Model: Matrix Factorization

- Map both users and items to a joint latent factor space
- item i : associated vector q_i
- user u : vector p_u
- $q_i^T p_u$ captures the interaction between u and i
- This approximates r_{ui} : user u 's rating of item i
- Factor rating matrix using SVD: obtain $, ,$
- Reduce the matrix to dimension
- Compute two resultant matrices:
- **Predicting task**



Latent Factor Models

- Goal: to find Q and P such that:

$$\min_{P,Q} \sum_{(i,u) \in R} (r_{ui} - q_i \cdot p_u)^2$$

		users			factors		
		items			users		
items	1	3	5	5	4		
		5	4	4	2	1	3
	2	4	1	2	3	4	3
	2	4	5	4		2	
	4	3	4	2		2	5
	1	3	3	2		4	

\approx

Q

		users										factors	
		PT											
1.1	-.2	.3	.5	-.2	-.5	.8	-.4	.3	1.4	2.4	-.9		
-.8	.7	.5	1.4	.3	-1	1.4	2.9	-.7	1.2	-.1	1.3		
2.1	-.4	.6	1.7	2.4	.9	-.3	.4	.8	.7	-.6	.1		
-.7	2.1	-2											
-1	7	.3											

Dealing with Missing Entries

- To address overfitting, introduce regularization:

$$\min_{P,Q} \underbrace{\sum_{\text{training}} (r_{ui} - q_i p_u)^2}_{\text{error}} + \left[\lambda_1 \sum_x \|p_u\|^2 + \lambda_2 \sum_x \|q_i\|^2 \right]$$

$\lambda_1, \lambda_2 \dots$ user set regularization parameters

1	3	4		
3	5		5	
4	5			5
3				
3				
2				
			?	
	2	1		?
3				
1				

Why use reviews?

- Reviews provide rich information.
- Reviews help to explain users' ratings.
- Reviews are useful at modeling new users and products: **one review tells us much more than one rating.**
- Reviews help us to discover the aspects of users' opinions
 - Modeling these factors improve recommendation
- Value for customer
 - Help explore space of options
 - Discover new things
- Value for company
 - Increase trust and loyalty
 - Increase conversion
 - Opportunities for promotion, persuasion
 - Obtain more knowledge about customers



The Value of Reviews

- Recommender systems played an important part in helping users find products.
 - Reviews were used by users but only relevant ones.
 - Any algorithmic recommender needs to communicate results to the user in a meaningful manner.
- The volume of reviews affects product conversion (upto 270% increase)
 - High priced items have a higher value for reviews.

Case Amazon: Ratings and Reviews as Part of Recommendations
Juha Leino and Kari-Jouko Räihä, RecSys 2007



The Value of Online Customer Reviews
Georgios Askalidis, Edward C. Malthouse
RecSys 2016



The importance of Explanations

- Explanations improve user engagement and satisfaction with a recommender system.
 - Controllability and transparency.
 - Why are these items recommended?
- Trust
 - user's confidence
 - Persuasive
 - Convince users to try/buy
 - Satisfaction
 - user's enjoyment
 - Help compose judgment
 - attributes to a rating

Reviews can help



Product recommendations enhanced with reviews



Part I

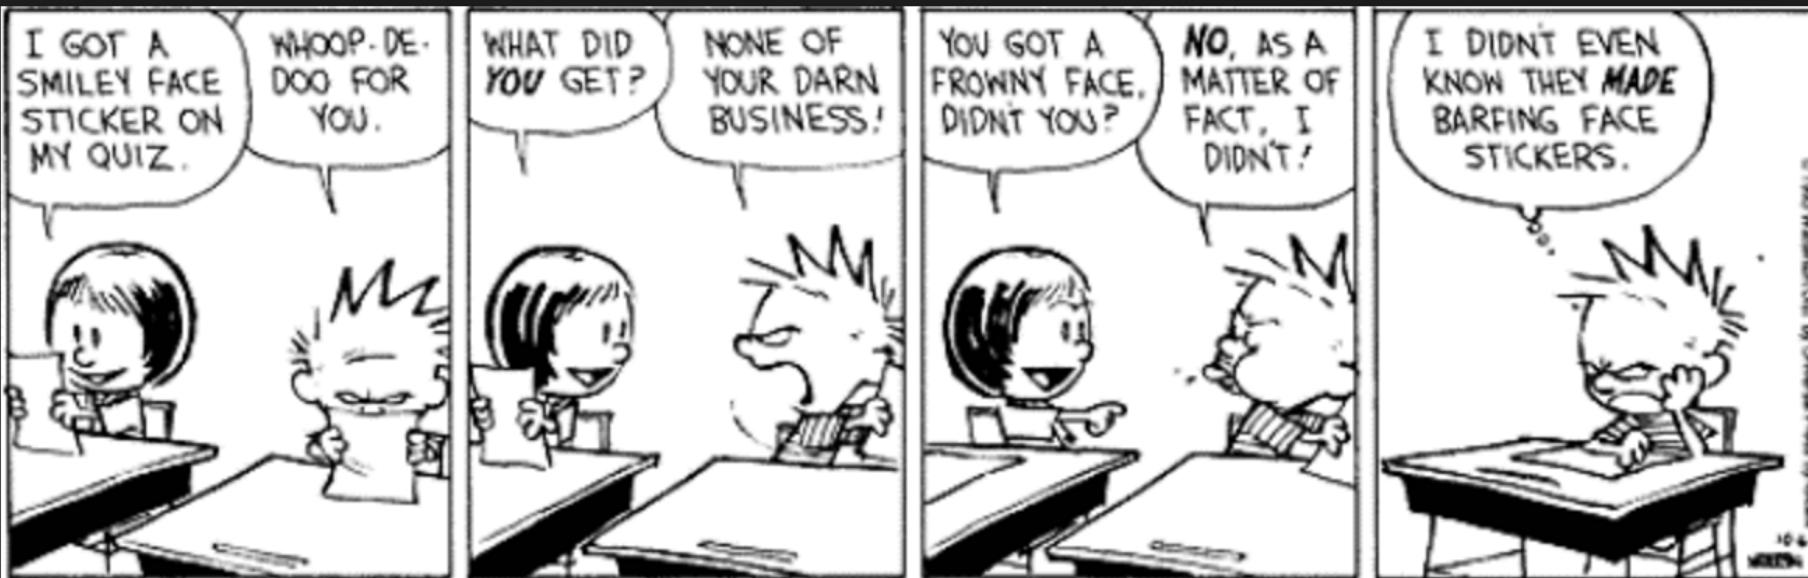
Muthusamy Chelliah

Director, Academic Engagement, Flipkart

Sudeshna Sarkar

Professor, CSE, IIT Kharagpur

Email: sudeshna@cse.iitkgp.ernet.in



Product Reviews



ADD TO CART

BUY NOW

Ratings and Reviews

4.4 ★

9,741 Ratings &
2,263 Reviews

5 ★

4 ★

3 ★

2 ★

1 ★

6,381

2,192

516

175

477

Most Helpful Positive Negative Most Recent By Certified Buyers

5 ★ Excellent

My first dslr....

Awesome package by flipkart....packed very nicely....and the gift received from offer....Motorola pulse headphone...amazing....
I m totally satisfied with both the products....thank you flipkart. ...

Rahul Prajapati Certified Buyer 7 Oct, 2016

1635

191

5 ★ Awesome

Awesome camera with stunning features. Love it

Ratikanta Sahoo Certified Buyer 4 Oct, 2016

792

141

1 ★ free motorola head set missing

Purchased the product on big billion day but product is not dispatched with all the offer as mentioned on the product page on big billion day .

Product specification



ADD TO CART

BUY NOW

Highlights

- 18 megapixel APS-C CMOS sensor & DIGIC 4+, WiFi and NFC supported, 9 point AF with 1 centre cross type AF point
- Effective Pixels: 18 MP
- Sensor Type: CMOS
- WiFi Available

Services

- 2 Years Canon India Warranty ?
- 10 Days Replacement Policy ?
- Cash on Delivery available ?

Seller

RetailNet

View 5 sellers starting from ₹23,499

Specifications

In The Box

EF-S18-55mm f/3.5-5.6 IS II EOS 1300D Body + Eyecup Ef + Body Cap Battery Charger LC-E10E Battery Pack LP-E10 Interface Cable Wide Strap EW-400D

General

Brand	Canon
Model Number	EOS 1300D
Model Name	Canon EOS 1300D
SLR Variant	(Body with EF-S 18 - 55 IS II)
Brand Color	Black
Type	DSLR
Color	Black
Effective Pixels	18 MP
Shooting Modes	Basic Zone Modes: Scene Intelligent Auto, Flash Off, Creative Auto, Portrait, Landscape, Close-up, Sports, Food, Night Portrait, Creative Zone Modes: Program AE, Shutter Priority AE, Aperture Priority AE, Manual exposure
Tripod Socket	Yes
Wifi	Yes
Face Detection	Yes

Definitions

- **Contributor:** The person or organization who is expressing their opinion in text.
- **Object:** An entity which can be a product, service, person, event, organization, or topic.
- **Review:** A contributor-generated text that contains their opinions about some **aspects** of the object.
- **Aspect:** The component, attribute or feature of the object that contributor has commented on.
- **Opinion:** An opinion on an aspect is a positive, neutral or negative view, attitude, emotion or appraisal on that aspect from a contributor.

Definitions

- **Contributor:** Contributor is the person or organization who is expressing the review.

5 of 5 people found the following review helpful
Very good camera overall, but with some issues for HD video, October 24, 2014
By T. Ishizue "Tom" (California) - See all my reviews
REAL NAME

- **Object:** An entity being observed by a person, even if it is a person, event, or thing.

This review is for the Canon PowerShot SX60 HS Digital Camera (Electronics)
I generally agree with the previous review comments that others have written, so I will not repeat them here. What I will add is my review of the video performance of this camera. First the strengths - The image stabilizer is simply phenomenal. I have not seen anything this good since Sony's Balanced Optical Image Stabilizer. What this means is that you can film literally on the run and the images will stay solid and not shaky (shakiness can make any video unusable). My test video is posted on YouTube titled, "Sunset at Fisherman's Wharf Using the Canon SX60 HS", where you can see how effective the image stabilizer is when I'm running around filming my 5-yr old daughter along the beach. The other strength is the ultra-wide 21mm (35mm equivalent) lens. This especially comes in handy for obtaining more stable shots than with a narrower focal length where every little movement becomes exponentially magnified.

- **Review:** Review is the opinion expressed about the object.

Now the weaknesses - I was really looking forward to utilizing Canon's first Powershot camera that offered an external mic input. When I connected my Rode Stereo VideoMic Pro, to my disappointment, the sound quality, while improved over the built-in microphones, was not as good as I expected. The sound was still somewhat hissy and did not have a full frequency range, especially at the low end. When I connected a lavalier mic to test my voice, it also was not clean as I expected. I checked the fullness of my other cameras when using the same microphones (Sony HDR-CX900, Sony HDR-MV1, and Panasonic GH3). After looking into this further, I read in the specifications that the audio recording codec was compressed Limited AAC, not uncompressed LPCM as was used in the SX40 and SX50. I contacted Canon Technical Support to submit this feedback as a possible area that can be remedied in a future firmware update (keeping my fingers crossed). Also, while not a major issue, Canon has historically not offered the ability to shoot stills while shooting video like some other manufacturers have.

- **Aspect:** Aspect is the specific feature or quality of the object that is being reviewed.

Overall, this is a very good camera, but for my use, if these minor issues can be addressed, I would think this would be an even better camera for run-and-gun HD video.

Help other customers find the most helpful reviews
Was this review helpful to you?

[Report abuse](#) | [Permalink](#)

[Comment](#)



Terminologies

- Features / Attributes / Aspect / Properties of a product or service
- Opinion
 - Overall opinion
 - Feature opinion
- Review helpfulness:
 - the number of “helpful” votes given by readers
 - Can be used to determine quality score

Aspect-based sentiment analysis (ABSA)

[Moghaddam '10]

Input

Canon GL2 Mini DVD Camcorder

... excellent zoom ... blurry lcd ... great picture quality ...
accurate zooming ... poor battery ... inaccurate screen ...
good quality ... affordable price ... poor display ...
inadequate battery life ... fantastic zoom ... great price ...

Output



Aspect	Head Terms	Rating
zoom	zoom, zooming	5
price	price	4
picture quality	picture quality, quality	4
battery life	battery life, battery	2
screen	screen, lcd, display	1
...		

Prediction

- [Liu '16]
- screen is clear and great
- **Aspect extraction:** screen
- **Opinion identification:** clear, great
- Polarity classification: clear is +ve
- Opinion separation/generality: clear (aspect-specific), great (general)
 - Understand consumer taste of products

FLIPKART CUSTOMER ASKS



Recommendation: User Needs

- **Help me narrow down on my choice**
 - I've shortlisted few cameras and need help deciding one
 - I'm almost leaning towards buying this speaker system but want to get further confidence and assurance on sound quality.

Flipkart Ratings-based Recommendations

- Product discovery with Rating as a signal
 - rank products with high quality ratings at the top
- Highlight Ratings to narrow down purchase
 - show ratings attributes that map to user intent

Similar products



Sony DSC-H300 Point & Shoot Camera

4.1★ (809)

₹14,698

Canon EOS 1300D DSLR Camera (Body with EF-S 18-55 IS II)

4.4★ (10,229)

Assured

₹26,490 ₹25,955 11% off

Nikon B700 Black Point & Shoot Camera

4.1★ (21)

₹22,490

[Cameras with good image stabilization >](#)
[Camera good for night mode >](#)

Top Rated phones with good Front camera

Because you searched for good Selfie phones



Samsung Galaxy S7 Edge (Gold Platinum, 32 GB)

4.4★ (1,042)

₹42,900

Moto X Play (With Turbo Charger) (Black, 32 GB)

4.2★ (51,602)

Assured

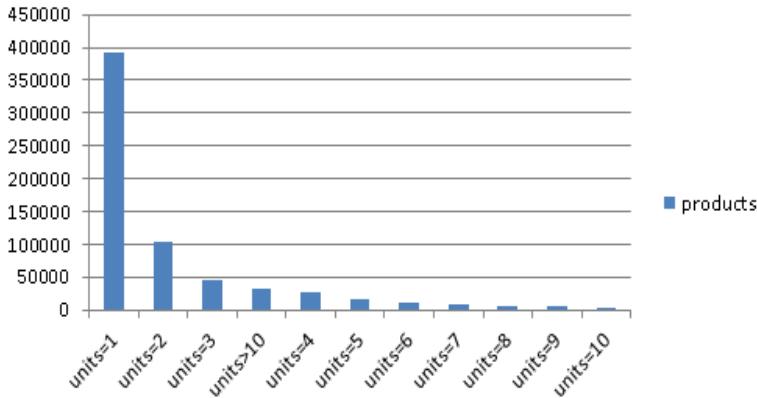
₹17,499 ₹18,999 7% off

Samsung Galaxy J7 (Gold, 16 GB)

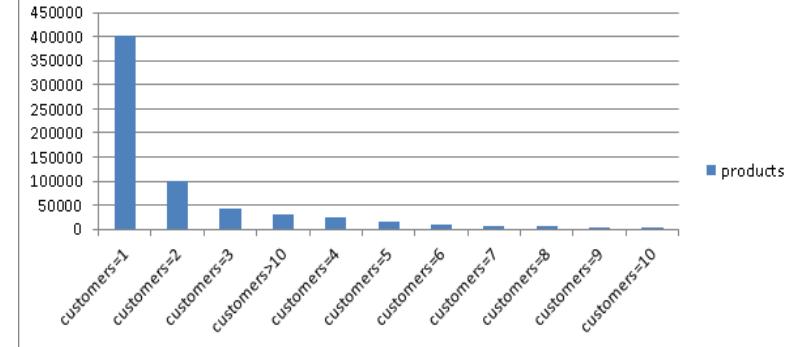
4.3★ (25,379)

₹12,490

Units bucket vs products



customer bucket vs products



Long tail (both in users/products):
induces sparsity/cold-start

Recommendations Leveraging Reviews

- Good coverage
- High quality and detailed reviews
- Parsing of reviews into meaningful clusters for Discovery
 - Picture Quality (4.5), Sound (3.8), Ease of Installation (2.5)

ABSA - User needs

- Individual preferences
 - easy way to shortlist products matching a criteria
 - compare two products on certain aspects
- Too many reviews to read
 - interested in only key elements of a product
 - be aware of product shortcomings before making the final buying decision

ABSA - Partner Needs

- Refine available selection per user preference
- Brands can validate future product design
- Correlation between aspect and sales is a good indicator of market demand



ABSA - agenda

- Statistics-based text mining (5 minutes)
- Topic models
 - Early (5 minutes)
 - Recent (10 minutes)
- Deep learning (10 minutes)

Research landscape – ABSA

Aspect extraction, Rating prediction

[Moghaddam '10]

Normalization [Guo '09][Bing '16]

Statistical text mining

ILDA [Moghaddam '11]
MG-LDA, JST, Maxent-LDA,
ASUM, CFACTS:overview
[Liu '11, Moghaddam '12]

Early topic models

HASM [Kim '13],
SpecLDA [Park '15]
LAST [Wang '16]

Recent topic models

Rating prediction

RecursiveNN [Socher'13]

Aspect extraction

CNN [Lu '14]

RecurrentNN [Liu '15]

LSTM/attention [Tang '16]

Deep learning

Rating prediction

1★ free motorola head set missing
Purchas ^ Back to top big billion day but product is not dispatched with all the offer as n on big billion day .
free motorola head set is missing , so giving one star .
parag dutta Certified Buyer 15 Oct, 2016

5★ Terrific purchase
Please make this price low seller I want to buy after my exam and for photography please pl
Flipkart Customer 22 Mar, 2017

5★ CANON Camara
Good product for Beginners like to learn photography and prompt delivery. Happy purchase.
.....
balasubramanian kuppusamy Certified Buyer 6 Jun, 2016

5★ Brilliant
owsam camera frndz go for it
Zaid Ashai Certified Buyer 21 Jan, 2017

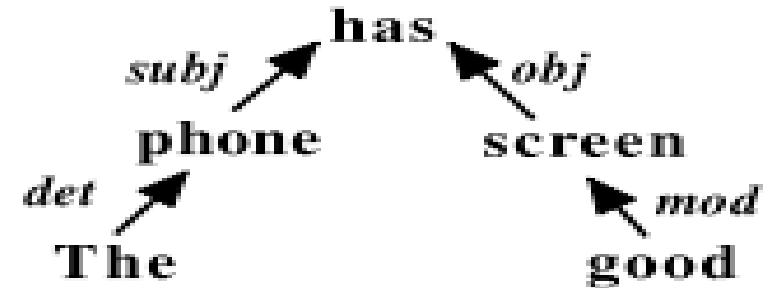
1★ Terrible product
Price is increasing changing everyday! Keep a correct price!
Flipkart Customer 15 Apr, 2017

[Moghaddam '10]

- Estimate score [1-5]
 - With kNN algorithm
 - Wordnet for similarity
- Aggregate rating
 - aspect-level estimate

Aspect extraction

- Frequency-based
 - Extract frequently occurring nouns based on seed words [Moghaddam '10]
- Pattern mining
 - Rules for identifying templates of POS-tagged tokens (e.g., _JJ_ASP)
- Clustering
 - Group aspect nouns [Guo '09]



Category	Product-feature Terms
Battery	<i>battery life; Lithium battery; AA Alkaline; AA batteries; battery charger; battery capability; battery pack; battery adapter; AA Lithium batteries; rechargeable battery;</i>
Memory	<i>memory; flash card; flash memory; memory card; memory capacity; SD card; sd memory; camera flash; digital memory; CF memory card;</i>

Attribute-aspect normalization

	Concept 1	Concept 2	Concept 3
	wireless, network, wi-fi, internet, connection	screen, resolution, display, brightness, monitor	battery, life, 6-cell, charge, capacity
Product 1	“draft-n wi-fi network,” “the integrated bluetooth connectivity”	“10.1 inches ultrawide display,” a superbright 10.1-inch lcd	“extended battery life,” “the traditional lithium battery”
Product 2	“fast ethernet connection,” “the bluetooth 2.0+edr technology”	“normal tft lcd,” “10.1 inches widescreen”	“65 hours battery life,” “a 6 cells battery”
Product 3	“ultimate wireless accessibility,” “wired ethernet lan”	“widescreen trubrite display,” “a 1024 × 600-pixel resolution”	“long battery life,” “6-cell lithium-ion battery”

The second row shows the top-five weighted terms in the popular features discovered from the customer reviews. Each cell of the table contains the popular attribute text fragments extracted from product description pages.

[Bing '16]

- Extract product attributes from description page
- Infer aspect popularity in reviews/map to feature
- Hidden CRF bridges vocabulary gap

References

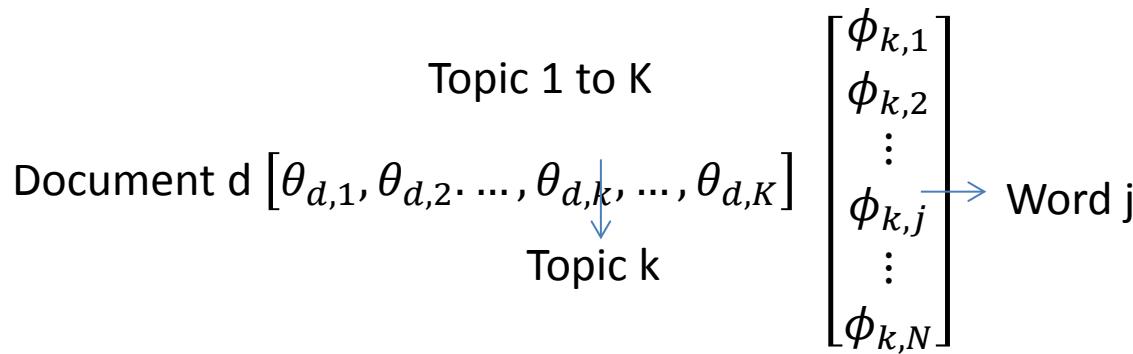
- Bing, L., Wong, T.L. and Lam, W. Unsupervised extraction of popular product attributes from e-commerce web sites by considering customer reviews. ACM TOIT '16.
- Guo, H., Zhu, H., Guo, Z., Zhang, X., & Su, Z. (2009, November). Product feature categorization with multilevel latent semantic association. CIKM 2009.
- Moghaddam, S., & Ester, M. Opinion digger: an unsupervised opinion miner from unstructured product reviews. CIKM 2010

ABSA - agenda

- Statistics-based text mining
- Topic models
 - Early (5 minutes)
 - Recent (10 minutes)
- Deep learning (10 minutes)

Topic Modeling: Latent Dirichlet Allocation (LDA)

- Given a document set D, assume for each , there is an associated K-dimensional topic distribution , which encodes the fraction of words in d that discuss each of the K topics. That is, words in document d discuss topic k with probability .
- Each topic k has an associated word distribution , which encodes the probability that a particular word is used for that topic k.
- and are random vectors each has a Dirichlet prior distribution.



Aspect-based sentiment analysis (ABSA)

[Moghaddam '10]

Input

Canon GL2 Mini DVD Camcorder

... excellent zoom ... blurry lcd ... great picture quality ...
accurate zooming ... poor battery ... inaccurate screen ...
good quality ... affordable price ... poor display ...
inadequate battery life ... fantastic zoom ... great price ...

Output



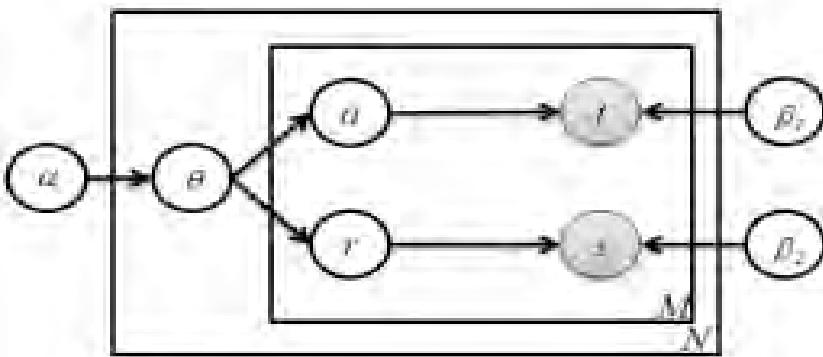
Aspect	Head Terms	Rating
zoom	zoom, zooming	5
price	price	4
picture quality	picture quality, quality	4
battery life	battery life, battery	2
screen	screen, lcd, display	1
...		

Prediction

- [Liu '16]
- screen is clear and great
- **Aspect extraction:** screen
- **Opinion identification:** clear, great
- Polarity classification: clear is +ve
- Opinion separation/generality: clear (aspect-specific), great (general)
 - Understand consumer taste of products

LDA – review mining

- Each item: finite mixture over an underlying set of latent variables
- Generative, probabilistic model for collection of discrete items
 - Review corpora generated first by choosing a value for aspect/rating pair $\langle a_m, r_m \rangle$ (θ)
 - Repeatedly sample M opinion phrases $\langle t_m, s_m \rangle$ conditioned on θ



Extract aspects/ratings together

Category	#Reviews	#Opinion Phrases	#Phrases per Product
Camcorder	197	1,694	211.77
Cel. Phone	630	7,642	955.23
Dig. Camera	707	10,435	1304.41
DVD Player	324	3,707	463.32
Mp3 Player	625	6,131	766.41
Overall	2,483	29,609	3701.14

- Topic modeling helps only with document clustering
- Independent aspect identification/ rating prediction leads to errors [Moghaddam '11]
 - Same sentiment word shows different opinions for various aspects
- Dataset
 - 29609 phrases, 2483 reviews, 5 categories from Epinions

Aspect-based sentiment analysis (ABSA)

[Moghaddam '10]

Input

Canon GL2 Mini DVD Camcorder

... excellent zoom ... blurry lcd ... great picture quality ...
accurate zooming ... poor battery ... inaccurate screen ...
good quality ... affordable price ... poor display ...
inadequate battery life ... fantastic zoom ... great price ...

Output

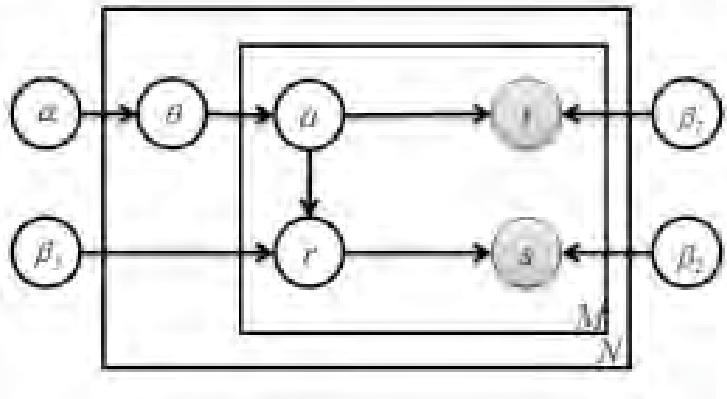


Aspect	Head Terms	Rating
zoom	zoom, zooming	5
price	price	4
picture quality	picture quality, quality	4
battery life	battery life, battery	2
screen	screen, lcd, display	1
...		

Prediction

- [Liu '16]
- screen is clear and great
- **Aspect extraction:** screen
- **Opinion identification:** clear, great
- Polarity classification: clear is +ve
- **Opinion separation/generality:** clear (aspect-specific), great (general)
 - Understand consumer taste of products

Interdependent LDA model



- [Moghaddam '11]
- Generate an aspect a_m from LDA
- Generate rating r_m conditioned on sample aspect
- Draw head term t_m and sentiment s_m conditioned on a_m and r_m respectively

References

Liu, B. Sentiment analysis. AAAI 2011.

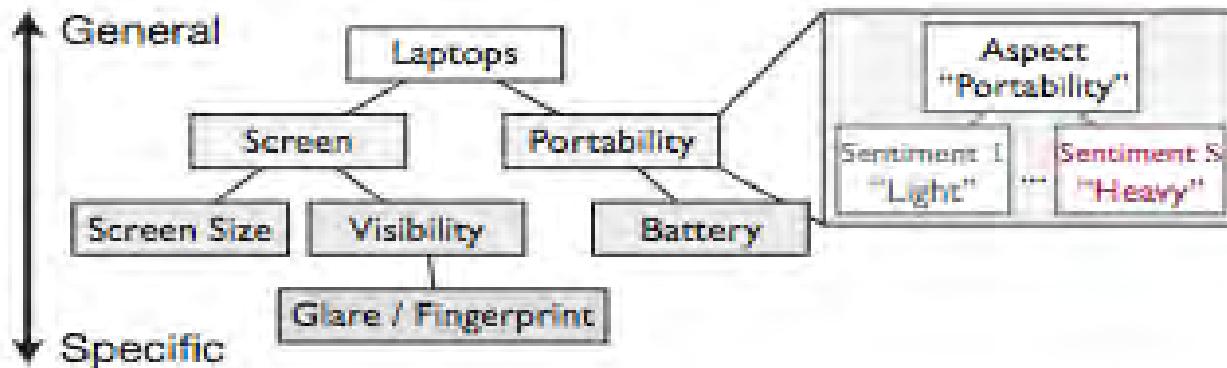
Moghaddam S, Ester M. ILDA: interdependent LDA model for learning latent aspects and their ratings from online product reviews. SIGIR '11

Moghaddam, S., & Ester, M. Aspect-based opinion mining from online reviews. SIGIR 2012.

Aspect-based sentiment analysis (ABSA)

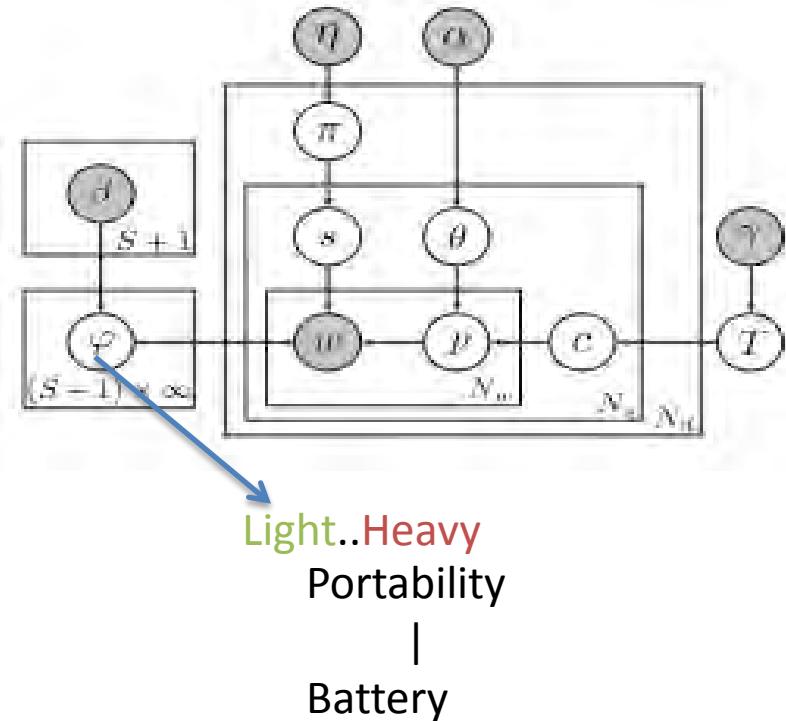
- Statistics-based text mining
- Topic models
 - Early
 - **Recent (10 minutes)**
- Deep learning (10 minutes)

Structure in Aspect-Sentiment



- Different consumers interested in different aspect granularities
- Polarity of general to aspect-specific sentiment words through sparse co-occurrence [Kim '13]
- Dataset
 - 10414 (laptop)/20862 (digital SLR) reviews, 825/260 targets: Amazon

Hierarchical aspect-sentiment model



- [Kim '13]

Jointly infer aspect-sentiment tree

- with a bayesian non-parametric model as prior

Aspect-sentiment node ϕ_k itself is a tree

- aspect topic at root and sentiment-polar (S) in leaves

Topics share semantic theme

- generated from dirichlet (beta) distribution

Sentiment-aspect hierarchy: Digital SLRs [Kim '13]

Digital SLRs			
(0) camera used lens can canon picture get take digital (+) lens picture great like can canon take good nikon (-) lens can canon picture digit mm shoot just set			
Image Quality	Lens	Operation	External Accessories
(0) camera image quality iso sensor (+) quality like high better sensor (-) picture quality size sensor raw	(0) lens mm camera kit zoom get (+) lens mm kit good af zoom great (-) mm camera kit buy ef purchase	(0) camera set mode can focus auto (+) set camera mode lcd view (-) camera set focus button mode	(0) camera card battery flash get (+) battery card raw software grip (-) card flash memory get battery
Image Sensor	Lens Performance	Menu Operation	Battery Grip
(0) sensor camera shake dust feature (+) dust sensor camera shake clean (-) work camera stabilize notshake	(0) lens mm vr camera kit lens used (+) lens kit mm good wide angle great (-) lens mm af dx nikon nikkor ed	(0) set mode camera can auto manual (+) mode set manual auto control (-) set mode manual menu change	(0) battery camera grip hand around (+) battery grip life hand comfort time (-) hand carry grip battery small bag
Low-Light Ambience	Lens Features	Live View Finder	Flash Memory
(0) iso noise camera image low light (+) quality image iso noise good high (-) iso noise low light high image set	(0) lens camera can use old system (+) lens stabilize body image range (-) lens af focus old camera older	(0) lcd button screen camera view set (+) lcd screen view live viewfind set (-) button shutter can press change	(0) card raw camera gb memory sd (+) raw card image file software photo (-) card memory gb sd cf flash speed

- [Kim '13]
- Polarity of sentiment words depend on aspect at various granularities

Aspect-based sentiment analysis (ABSA)

[Moghaddam '10]

Input

Canon GL2 Mini DVD Camcorder

... excellent zoom ... blurry lcd ... great picture quality ...
accurate zooming ... poor battery ... inaccurate screen ...
good quality ... affordable price ... poor display ...
inadequate battery life ... fantastic zoom ... great price ...

Output



Aspect	Head Terms	Rating
zoom	zoom, zooming	5
price	price	4
picture quality	picture quality, quality	4
battery life	battery life, battery	2
screen	screen, lcd, display	1
...

- [Liu '16]
- screen is clear and great
- **Aspect extraction:** screen
- **Opinion identification:** clear, great
- **Polarity classification:** clear is +ve
- **Opinion separation/generality:** clear (aspect-specific), great (general)
 - Understand consumer taste of products



Prediction



Aspect-sentiment topic

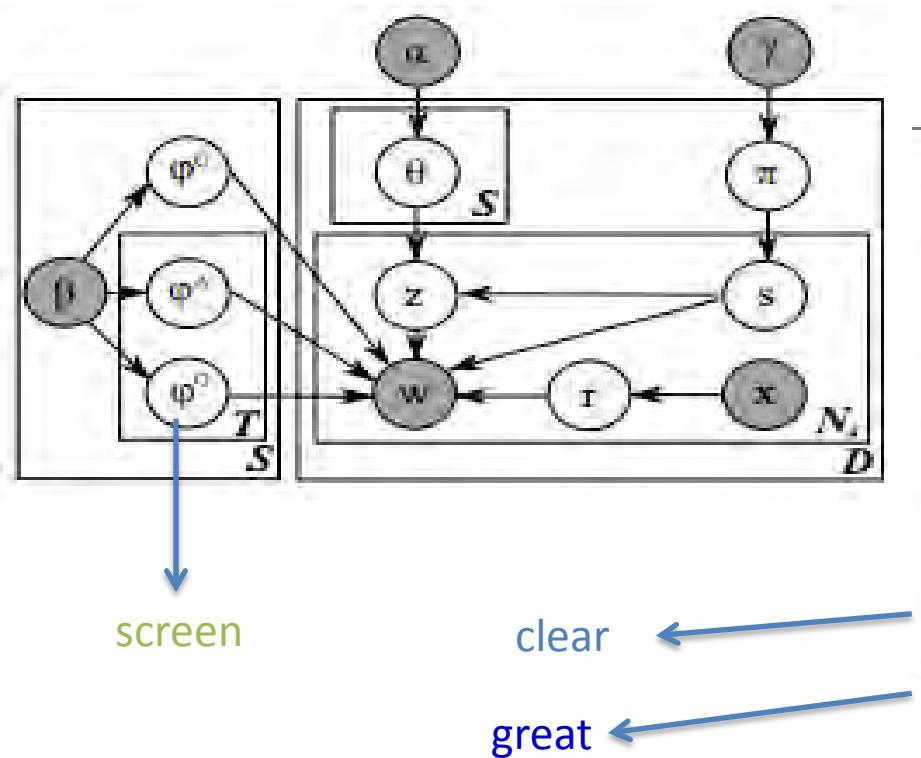
- [Wang '16]
- Generated topics inconsistent with human judgment
 - Wrong identification of opinion words (**nice** screen)
- Models based on co-occurrence suffer
 - Opinions (**smooth** screen) NOT identified if mixed in same topic

Lifelong learning across categories

- [Wang '16]
- Extract/accumulate knowledge from past
- Use discovered knowledge for future learning
 - Like human experience
- Dataset:
 - 50 product domains with 1000 reviews each - Amazon



JAST Model



[Wang '16]

S	the number of sentiment polarities
D	the number of documents
T	the number of aspect topics
V	the number of words or terms in vocabulary
N_d	the number of words in document d
s, d, z	sentiment polarity, document, topic
w, x, r	word, lexicon indicator, word type
π	multinomial distribution over sentiments
θ	multinomial distribution over topics or aspects
φ^G	multinomial distribution over general opinion words
φ^A	multinomial distribution over aspect words
φ^C	multinomial distribution over aspect-specific opinion words
α, β, γ	Dirichlet prior for θ, φ, π

Topic quality evaluation: opinion

Battery (Negative)			Shipping&Order (Positive)		
LAST	JAST	ASUM	LAST	JAST	ASUM
die	old	<i>problem</i>	new	free	<i>great</i>
dead	die	hot	free	<i>happy</i>	<i>good</i>
short	fail	<i>bad</i>	fast	fast	quickly
drain	<i>suck</i>	die	quick	<i>pleased</i>	<i>well</i>
fail	useless	<i>original</i>	refund	refund	<i>love</i>
old	hassle	old	promptly	<i>recommend</i>	<i>perfect</i>
hassle	<i>bad</i>	<i>new</i>	original	new	<i>nice</i>
<i>wrong</i>	<i>concern</i>	<i>long</i>	correct	works	<i>perfectly</i>
useless	bother	break	works	quick	new
<i>complain</i>	<i>nervous</i>	<i>hate</i>	accurate	promptly	fast

- [Wang '16]
- Incorrect (aspect-specific/polarity) (new, original)
- Incorrect (aspect-specific) (bad, suck)

Topic quality evaluation: aspect

Battery			Shipping&Order		
LAST	JAST	ASUM	LAST	JAST	ASUM
battery	battery	charge	order	arrive	<i>screen</i>
charge	charge	battery	receive	receive	receive
hour	life	recharge	arrive	order	arrive
life	hour	<i>iphone</i>	shipping	purchase	order
power	<i>device</i>	<i>sd</i>	ship	<i>expect</i>	<i>privacy</i>
charger	cable	<i>card</i>	today	send	cost
recharge	<i>phone</i>	receive	delivery	ship	money
<i>night</i>	<i>ipad</i>	replacement	usual	shipping	<i>monitor</i>
outlet	power	<i>purcharse</i>	<i>expect</i>	back	purchase
aaa	plug	<i>star</i>	<i>manner</i>	<i>seller</i>	<i>seller</i>

- [Wang '16]
- Joint modeling improves aspect quality while mining coherent opinions

Attribute-aspect normalization

	Concept 1	Concept 2	Concept 3
	wireless, network, wi-fi, internet, connection	screen, resolution, display, brightness, monitor	battery, life, 6-cell, charge, capacity
Product 1	“draft-n wi-fi network,” “the integrated bluetooth connectivity”	“10.1 inches ultrawide display,” a superbright 10.1-inch lcd	“extended battery life,” “the traditional lithium battery”
Product 2	“fast ethernet connection,” “the bluetooth 2.0+edr technology”	“normal tft lcd,” “10.1 inches widescreen”	“65 hours battery life,” “a 6 cells battery”
Product 3	“ultimate wireless accessibility,” “wired ethernet lan”	“widescreen trubrite display,” “a 1024 × 600-pixel resolution”	“long battery life,” “6-cell lithium-ion battery”

The second row shows the top-five weighted terms in the popular features discovered from the customer reviews. Each cell of the table contains the popular attribute text fragments extracted from product description pages.

[Bing '16]

- Extract product attributes from description page
- Infer aspect popularity in reviews/map to feature
- Hidden CRF bridges vocabulary gap

Augmented specification: camera

Importance	Feature	Value
1	AE/AF Control	FlexZone
2	Face Detection	Automatic Face Tracking technology
3	Digital Video Format	MOV
4	Image Recording Format	
5	Max Video Resolution	
6	AV Interface	
7	Manufacturer	

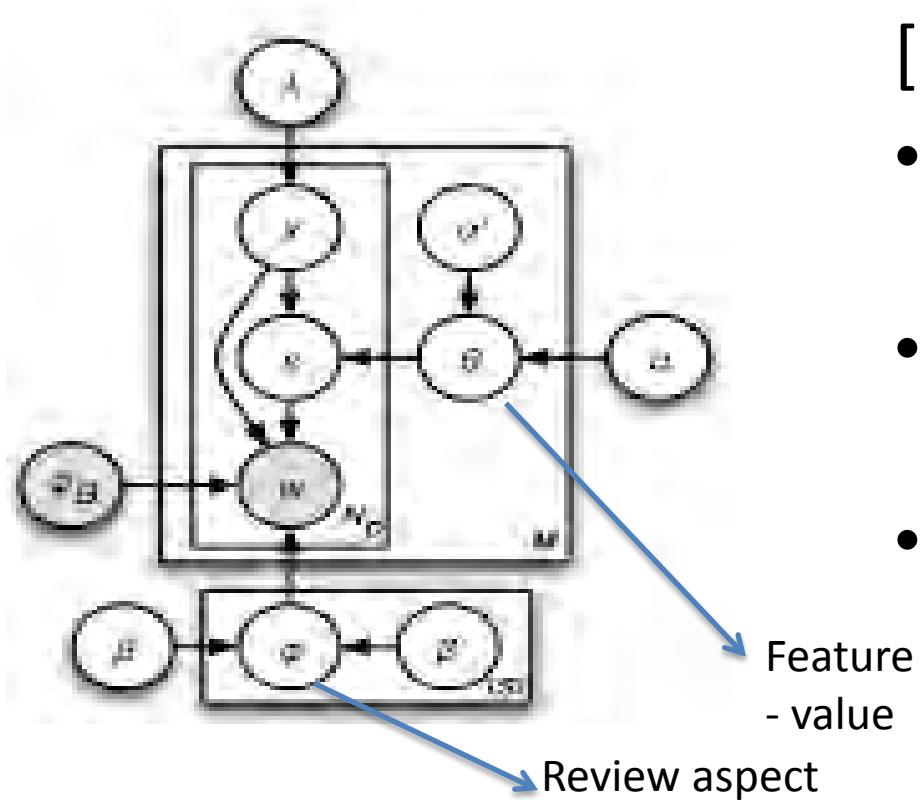
Product-specific Words **70d screen lcd**

MOV

+ Movie mode is .MOV which is not the standard mpg
+ Only fault is that you can only save video in .mov format.
+ Without spending some dough on a movie editor that supports .mov I can't edit any of the pictures with some of the commonly available free video editors such as Windows Movie Editor.

- [Park '15]
- Advanced features hard for novice customers
- Specs can be enhanced with reviews
 - Opinions on feature value/importance
 - Product-specific words

DuanLDA



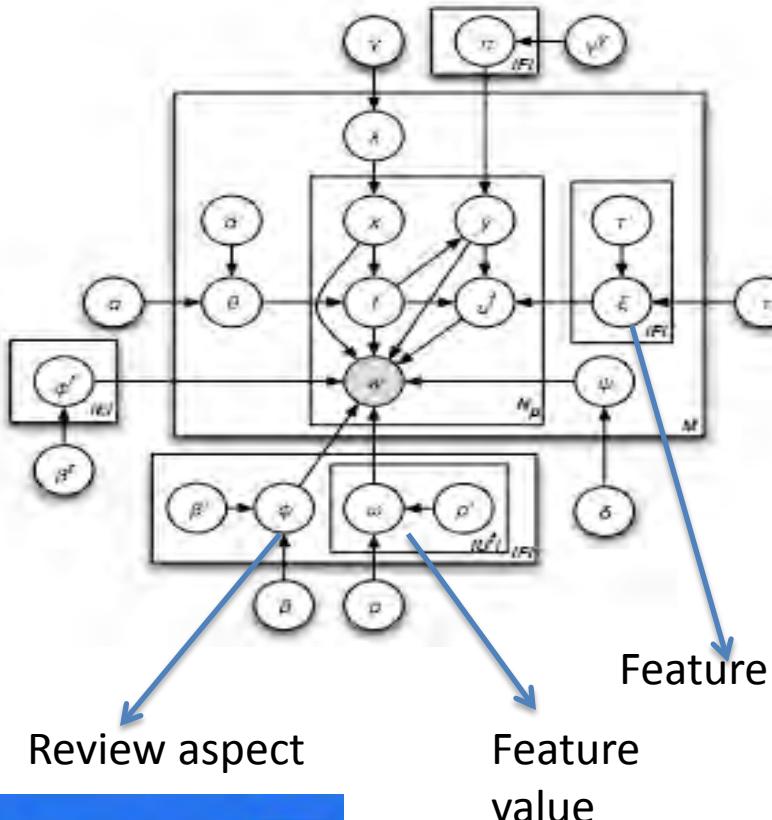
[Park '15]

- Each document is concatenated review
- s is feature-value pair/specification topic
- Switch x decides background/spec. topic

SpecLDA

[Park '15]

- Feature-value pairs with same feature share info
- Organized as hierarchy for better topic estimation
- Capture user preference to indicate a pair
 - for feature-/value-related words
- Employ value-word prior in addition to feature-word



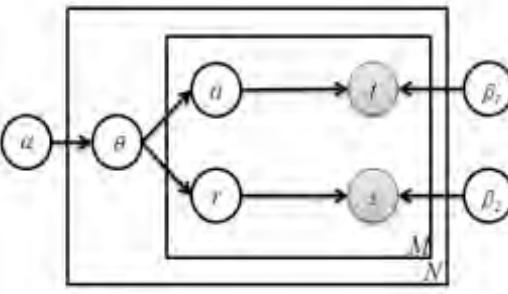
Evaluation: top words for a spec. topic

DuanLDA	SpecLDA-	SpecLDA-V	SpecLDA
display	digital	lcd	lcd
type	led	display	2.5
phtography	2.5	2.5	display
font	display	screen	screen
fluorescent	huge	type	large
informative	technology	large	inch
channel	expensive	inch	monitor
triple	type	monitor	articulate
resultant	large	photography	1.8
lcd	outstanding	font	230,000
printout	icon	crack	panel
a10	silver	symbol	bright
colorful	follow	1.8	icon
information	rubber	icon	rotatable
info	photography	bright	sunlight
horizontal	salesman	brightness	brightness
picture/video	font	articulate	viewfinder
infinite	info	salesman	tilt
2aa	blessing	range	viewer
3.0.	symbol	informative	flippable

[Park '15]

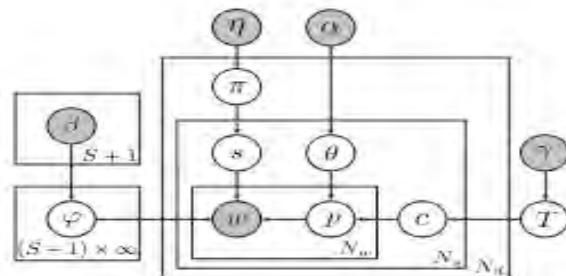
- Spec. (“Display – type”, “2.5in LCD display”)
- DuanLDA – does not use value word prior
- SpecLDA – feature and value topics form a cluster

ABSA Summary (Topic models)

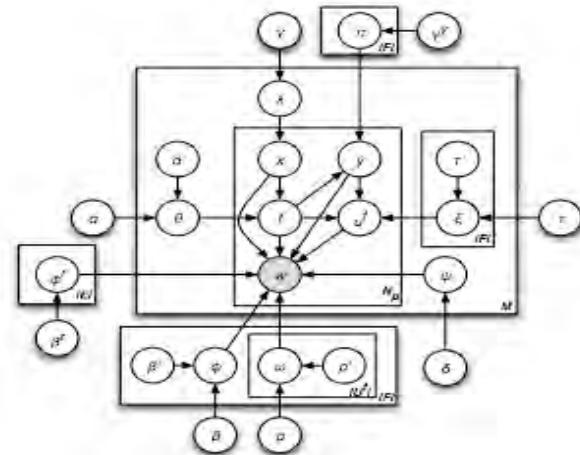


LDA: sample opinion phrases, not all terms

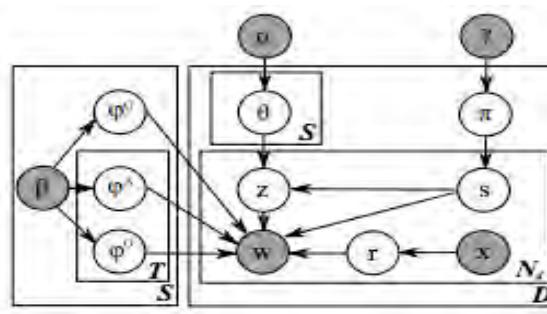
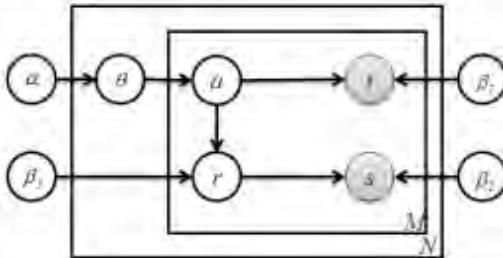
ILDA: Correspondence betn. aspect/rating



HASM: aspect-sentiment tree at each with topic root and S sentiment leaves



SpecLDA: feature-, value- word, aspect



JAST: aspect: (term, opinion), polarity, generic opinion

References

Kim, S., Zhang, J., Chen, Z., Oh, A. H., & Liu, S. A Hierarchical Aspect-Sentiment Model for Online Reviews. AAAI 2013.

Park, D.H., Zhai, C. and Guo, L. Speclda: Modeling product reviews and specifications to generate augmented specifications. SDM 2015.

Wang, S., Chen, Z., & Liu, B. Mining aspect-specific opinion using a holistic lifelong topic model. WWW 2016.



Aspect-based sentiment analysis (ABSA)

- Statistics-based text mining
- Topic models
 - Early
 - Recent
- Deep learning (10 minutes)

Aspect-based sentiment analysis (ABSA)

[Moghaddam '10]

Input

Canon GL2 Mini DVD Camcorder

... excellent zoom ... blurry lcd ... great picture quality ...
accurate zooming ... poor battery ... inaccurate screen ...
good quality ... affordable price ... poor display ...
inadequate battery life ... fantastic zoom ... great price ...

Output



Aspect	Head Terms	Rating
zoom	zoom, zooming	5
price	price	4
picture quality	picture quality, quality	4
battery life	battery life, battery	2
screen	screen, lcd, display	1
...		

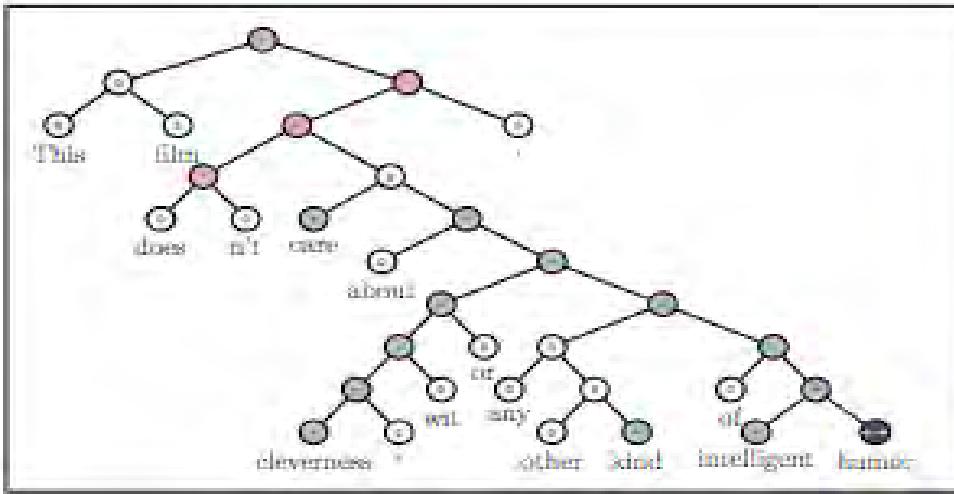
Prediction

- [Liu '16]
- screen is clear and great
- **Aspect extraction:** screen
- **Opinion identification:** clear, great
- Polarity classification: clear is +ve
- Opinion separation/generality: clear (aspect-specific), great (general)
 - Understand consumer taste of products

ABSA/Deep learning: Overview

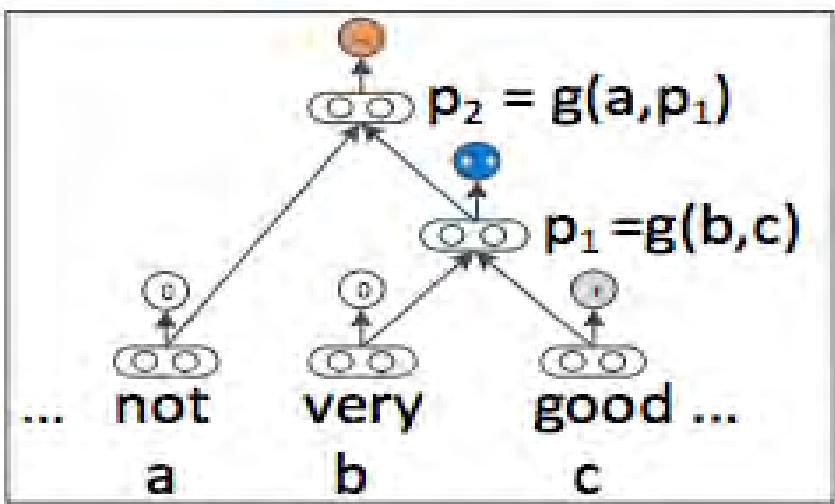
- Rating prediction
 - Recursive neural networks
- Aspect extraction
 - CNN
 - Recurrent neural networks
 - LSTM/attention

Semantic compositionality



- Rating prediction [Socher '13]
- Word spaces cannot express meaning of longer phrases in a principled way
- Sentiment detection requires powerful models/supervised resources

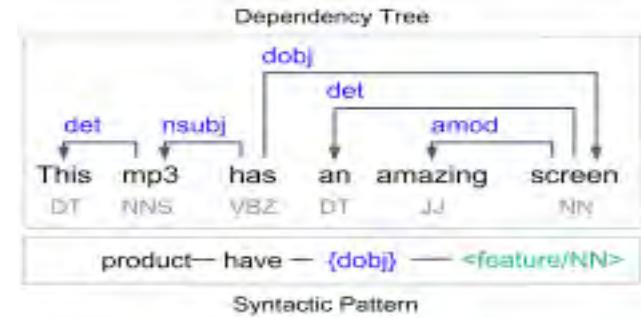
Recursive deep models



- Rating prediction [Socher '13]
- Parent vectors computed bottom-up
- compositionality function varies for different models
- Node vectors used as features for classifier

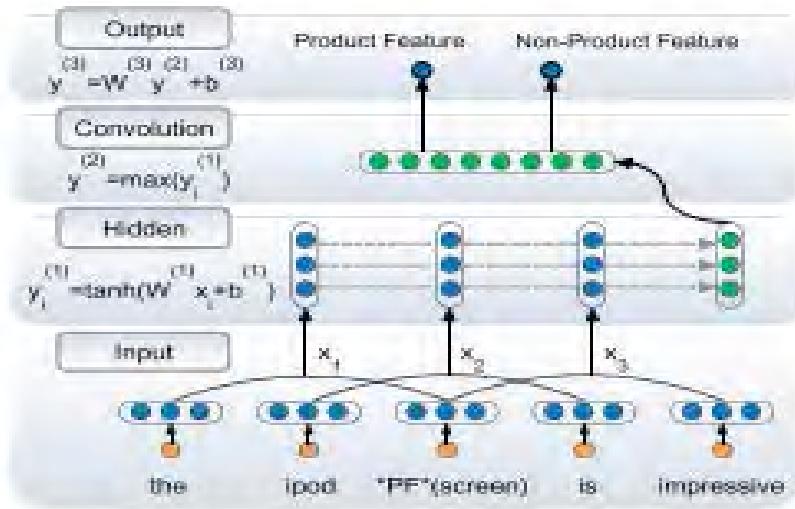
Semantic product feature mining

- Aspect extraction [Xu '14]
- Reviews features mined using syntactic patterns
- Syntax helps leverage only context
 - Suffering from sparsity of discrete information
- Term-product association needs semantic clues



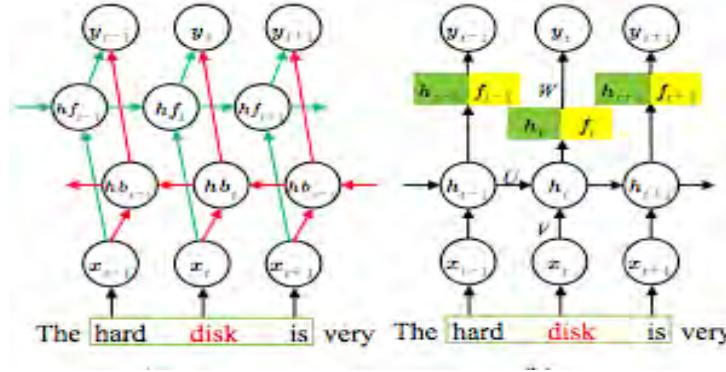
- a) *This player has an fm tuner.*
- b) *This mp3 supports wma file.*
- c) *This review has helped people a lot.*
- d) *This mp3 has some flaws.*

Convolutional neural network



- Aspect extraction [Xu '14]
- Similarity graph
 - Encodes lexical clue
- CNN
 - Captures context
- Label propagation
 - Captures both semantic clues

Recurrent neural networks

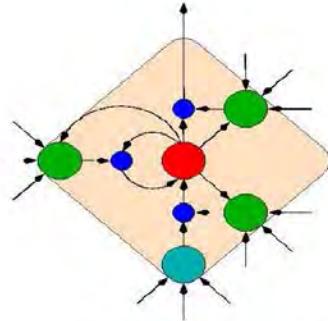


- Aspect extraction [Liu '15]
- window-based approach for short-term dependencies
- bidirectional links for long-range dependencies
- linguistic features to guide training

Deep memory network: LSTM

LSTM cell (current standard)

- Red: linear unit, self-weight 1.0 - the error carousel
- Green: sigmoid gates open / protect access to error flow; forget gate (left) resets
- Blue: multiplications



1/14/2003

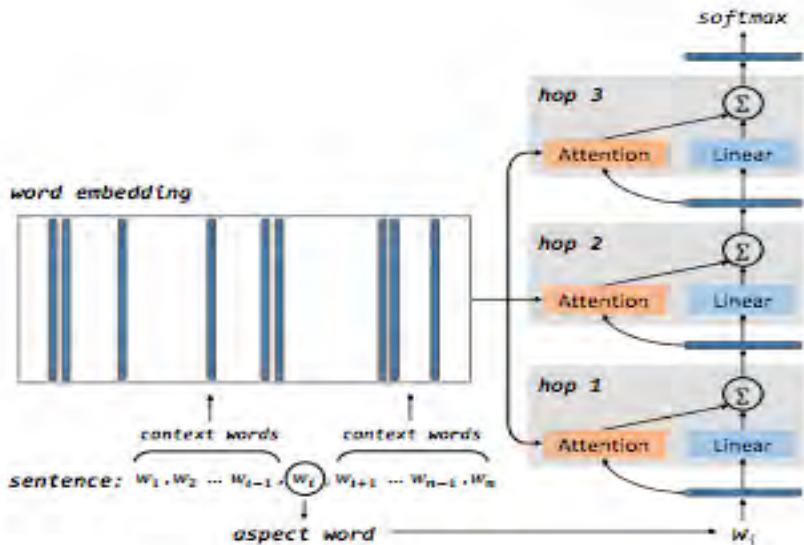
Juergen Schmidhuber (IDSIA)

LSTM tutorial

16

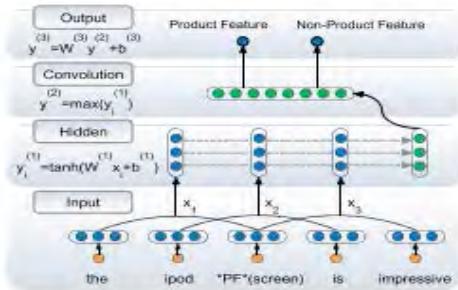
- Aspect extraction [Tang '16]
- *Fast processor but small memory*
- Need to capture semantic relations of aspect/context words
- LSTM incapable of exhibiting context clues
- Selectively focus instead on context parts

Deep memory network: attention



- Aspect extraction [Tang '16]
- Multiple hops are stacked
 - so abstractive evidence is selected from external memory
- Attention mechanism
 - weight to lower position while computing upper representation
- linear/attention layers with shared parameters in each hop
- Context word importance based on distance to aspect

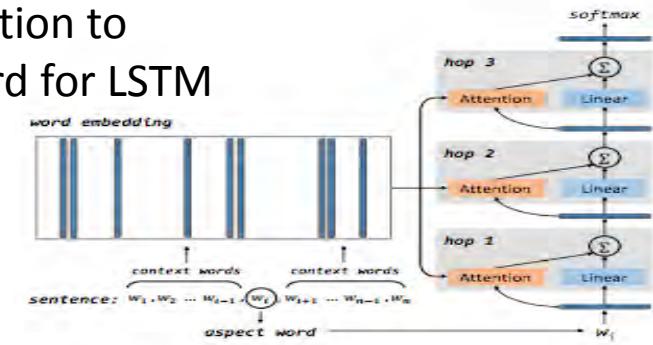
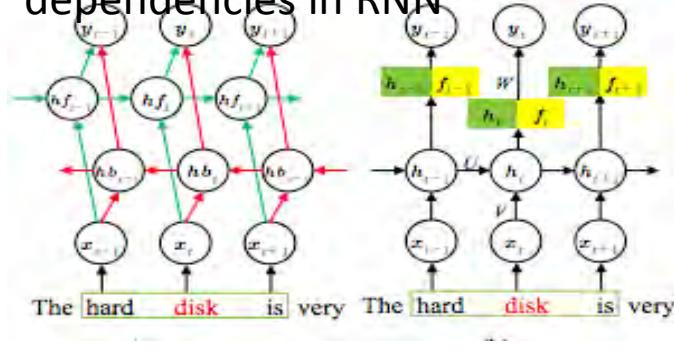
ABSA summary (deep learning)



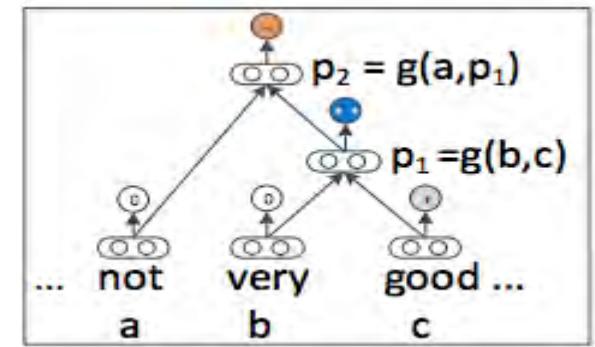
Content/location attention to emphasize context word for LSTM

Similarity graph for lexical and CNN for context semantics

Window for short-term and bidirectionality for long-range dependencies in RNN



Parent vectors computed bottom-up in RecNN for semantic compositionality



References

Liu, P., Joty, S. R., & Meng, H. M. Fine-grained Opinion Mining with Recurrent Neural Networks and Word Embeddings. EMNLP '15.

Socher, R., Perelygin, A., Wu, J.Y., Chuang, J., Manning, C.D., Ng, A.Y. and Potts, C., Recursive deep models for semantic compositionality over a sentiment treebank, EMNLP '13.

Tang, D., Qin, B. and Liu, T. Aspect level sentiment classification with deep memory network. EMNLP '16

Xu, L., Liu, K., Lai, S. and Zhao, J.. Product Feature Mining: Semantic Clues versus Syntactic Constituents. ACL 2014.

FLIPKART ABSA



Trusted advisor

Objective

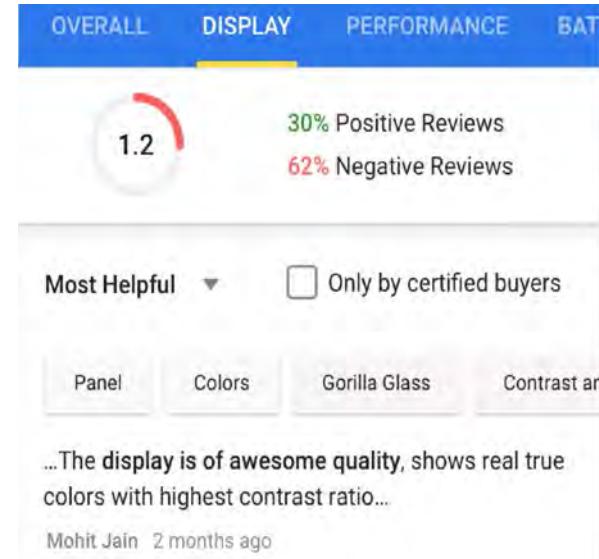
- Transform Flipkart into a research destination
- Shorten the purchase cycle for a user

Aspect Reviews

- Identify what users care about while buying a specific product
- Organize/classify reviews into these aspects & sub-aspects

Aspect Ratings

- Organize sentiments (polarity & strength)
- Summarize sentiments in the form of a user rating at an aspect level

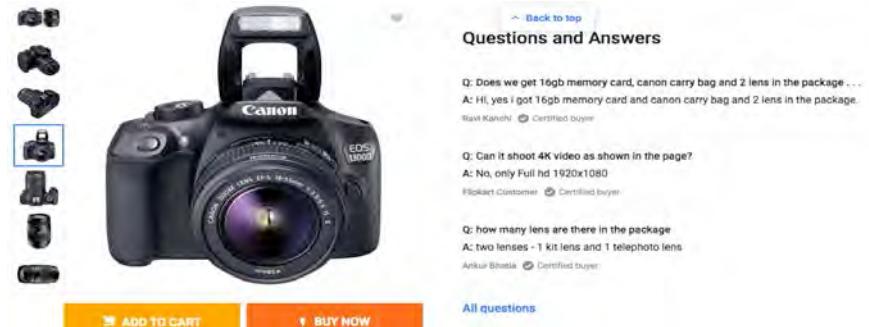


योग: कर्मसु कौशलम्

Future directions

- Product comparison
 - Model entities assessed directly in a single review sentence [Sikchi '16]
 - Leverage product spec. to learn attribute importance
- Question/answer
 - Summarize reviews with real-world questions
 - Rerank initial candidates for diversity [Liu '16]

Aspects	Camcorder1	Camcorder2
zoom	4	2
sound	1	3
screen	2	4
price	3	3
size	4	2
Overall Rating	3	3



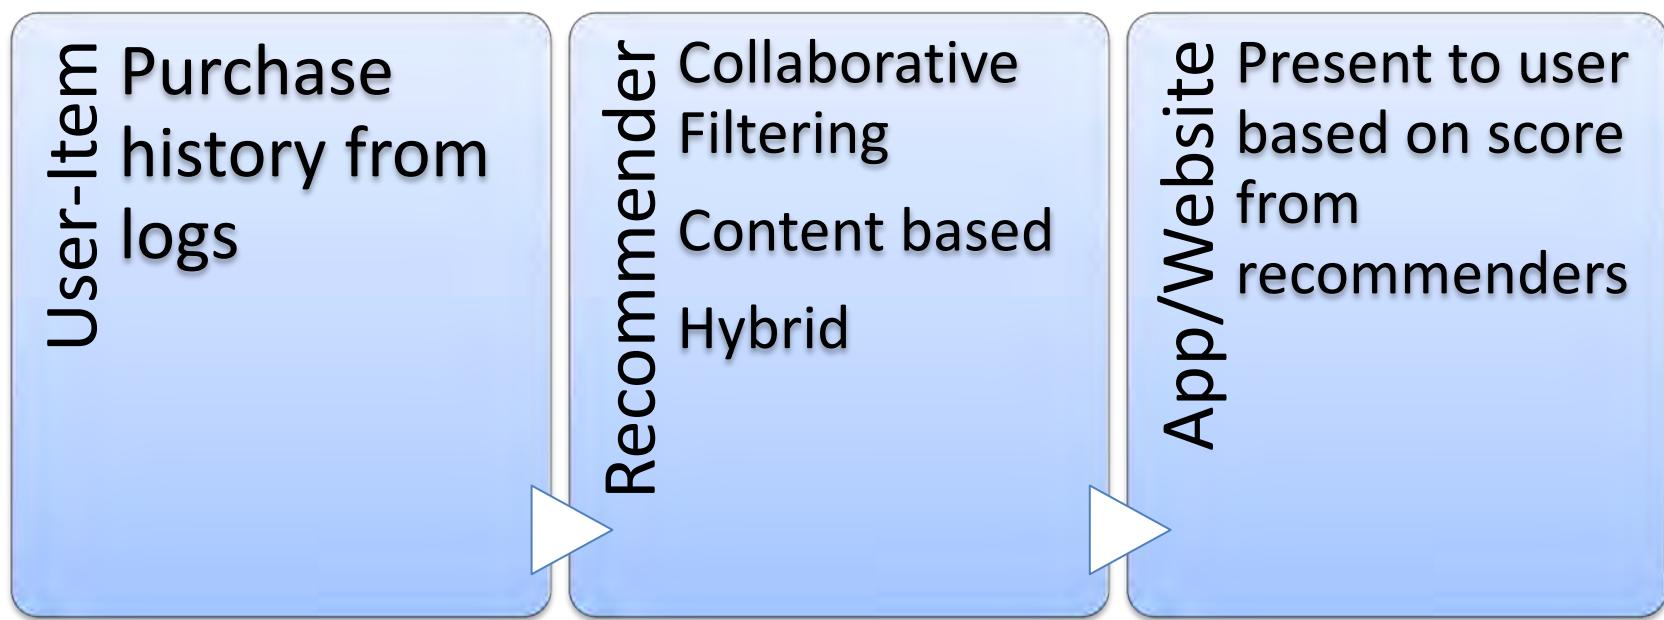
References

- Liu, M., Fang, Y., Park, D.H., Hu, X. and Yu, Z. Retrieving Non-Redundant Questions to Summarize a Product Review. SIGIR 2016.
- Sikchi, A., Goyal, P. and Datta, S., Peq: An explainable, specification-based, aspect-oriented product comparator for e-commerce. CIKM 2016.

FLIPKART RECOMMENDATION



Traditional View



Handling Cold-Start

- L1 ranking – leverage content (e-commerce has structured data), business provided attributes (CTR, conversion rate)
- L2 ranking - many online learning algorithms – AdaGrad, FTRL, AdPredictor for real-time learning (per-coordinate learning) , can handle new users/products very soon
- Other sources: Reviews, ratings
 - Some issues
 - Reviews/ratings are always biased to extremes
 - Ratings: Users very satisfied or unsatisfied
 - Reviews: more information to type. Biased towards angst/unsatisfactory experience

Future directions

- Bundle [Pathak '17]
 - Personalized/not considered in isolation
 - Assess size/compatibility
- Upsell [Lu '14]
 - Maximize revenue thru utility, saturation and competing products

Frequently bought together



Total price: ₹25,369.00

Add both to Cart

This item: Canon EOS 1300D 18MP Digital SLR Camera (Black) with 18-55mm IS II Lens, 16GB Card and Carry Case ₹24,990.00

Original Scratchgard Ultra Clear Screen Protector for Canon EOS 1300D ₹379.00

Similar products



Sony DSC-H300 Point & Shoot Camera

4.1 ★ (827) ₹14,490



Canon EOS 1300D DSLR Camera (Body with EF-S 18 - 55 IS II)

4.4 ★ (10,344) ₹24,990 ₹29,995 16% off



Nikon B700 Black Point & Shoot Camera

4.1 ★ (21) ₹23,950

References

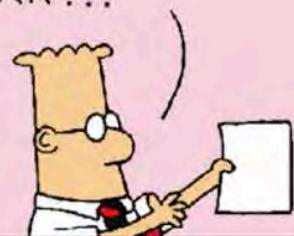
- Pathak, A., Gupta, K., & McAuley, J. Generating and Personalizing Bundle Recommendations on Steam. SIGIR '17.
- Lu, W., Chen, S., Li, K., & Lakshmanan, L. V. Show me the money: Dynamic recommendations for revenue maximization, VLDB '14.

SO, YOU IGNORED MY RECOMMENDATION AND BOUGHT A LOW-COST SYSTEM THAT'S TOTALLY INADEQUATE...



S. ADAMS E-mail: SCOTTADAMS@AOL.COM

YOU COMPENSATED FOR THIS BLUNDER BY MAKING IT PART OF MY OBJECTIVES TO MAKE THE SYSTEM WORK ...



YOU'LL GET A BONUS FOR SAVING MONEY. I'LL GET FIRED, THUS SAVING MORE MONEY AND EARNING YOU ANOTHER BONUS.



12/7 © 1995 United Feature Syndicate, Inc.(NYC)



Product recommendations enhanced with reviews



Part 2

Muthusamy Chelliah

Director, Academic Engagement, Flipkart

Sudeshna Sarkar

Professor, IIT Kharagpur

Email: sudeshn@cse.iitkgp.ernet.in

Review Enhanced Recommendation

1. Introduction
2. Review mining
 - a) Background: Features and Sentiment, Latent features
 - b) Statistical approaches. LDA
 - c) Recent topic models
 - d) Deep Learning (ABSA)
 - e) Flipkart ABSA/Recommendation
3. Review based Recommendation
 - a) Handling Data Sparsity
 - b) Topic Model, Matrix Factorization and Mixture Model based Recommendation
 - c) Deep learning based Recommendation
4. Explainable Product Recommendations
5. Summary, Future Trend

Handling Data Sparsity

Explainable Recommendation

REVIEW AWARE RECOMMENDATION



Issues with Recommender Systems

Collaborative Filtering

- Cold Start
 - New items, New Users
- Sparsity
 - it is hard to find users who rated the same items.
- Popularity Bias
 - Cannot recommend items to users with unique tastes.

Content Based

- Hard to identify quality judgement.
- Ignores user's ratings.

Review Based

- Text information from reviews address sparseness of rating matrix.
 - Bag of words model
- More sophisticated models required to capture contextual information in documents for deeper understanding: CNN, etc.
- Enable Explanations

Review Aware Recommender Systems

Traditional Models

Weight on quality of ratings, Recsys'12

Use reviews as content, Recsys'13

Aspect Weighting, WI'15

TriRank, CIKM'15

Topic Models and Latent Factors

HFT, Recsys '13

RMR, Recsys'14

CMR, CIKM'14

TopicMF, AAAI'14

FLAME, WSDM'15

Joint Models of Aspects and Ratings

JMARS, KDD'14

EFM, SIGIR'14

TPCF, IJCAI'13

SULM, KDD'17

Deep Learning Models

ConvMF, Recsys'16

DeepCONN, WSDM'17

TransNets, Recsys'17

D-Attn, Recsys'17

Using weights based on reviews

- Attach weights or quality scores to the ratings.
- The quality scores are determined from the corresponding review.
- Incorporate quality scores as weights for the ratings in the basic PMF framework
- Propose a unified framework for performing the three tasks:
 - aspect identification
 - aspect-based rating inference
 - weight estimation.

Review Quality Aware Collaborative Filtering, by Sindhu Raghavan, Suriya Gunasekar, Joydeep Ghosh, Recsys'12

Review Mining for Estimating Users' Ratings and Weights for Product Aspects, by Feng Wang and Li Chen, WI'15

Sentimental Product Recommendation

1. Convert unstructured reviews into rich product descriptions
 - a) Extract Review Features using shallow NLP
 - b) Evaluate Feature Sentiment by opinion pattern mining
2. Content-based recommendation approach based on feature similarity to a query product
 - a) Similarity-Based Recommendation
 - b) Sentiment-Enhanced Recommendation: seek products with better sentiment than the query product.

Sentimental Product Recommendation, by Ruihai Dong, Michael P. O'Mahony, Markus Schaal, Kevin McCarthy, Barry Smyth, Recsys 2013



Review Aware Recommender Systems

Topic Models and Latent Factors

HFT, McAuley & Leskovec, Recsys 13,
LDA+MF

RMR, Ling et al, Recsys'14,
LDA+PMF

CMR, CIKM'14, Xu et al,
LDA+PMF+User factors

TopicMF, Baoetal, AAAI'14
NMF+MF

Joint Models of Aspects and Ratings

JMARS, Diao et al, KDD'14,
Graphical Models

EFM, Zhang et al, SIGIR'14,
Collective NMF

TPCF, Musat et al, IJCAI'13
User topical profiles

SULM, KDD'17

HFT: Linking Reviews

1. Latent Dirichlet Allocation (LDA) finds low dimensional structure in review text (topic representation)
2. SVD learns latent factors for each item.
 - Statistical models combine latent dimensions in rating data with topics in review text.
 - Use a transform that aligns latent rating and review terms.
 - Link them using an objective that combines the accuracy of rating prediction (in terms of MSE) with the likelihood of the review corpus (using a topic model).

Hidden Factors and Hidden Topics: Understanding Rating Dimensions with Review Text: McAuley and Leskovec, Recsys 2013



HFT: Optimization

- A standard recommender system : $\hat{r}_{u,i} = p_u \cdot q_i$
- Minimize MSE:

$$\min_{P,Q} \frac{1}{|\mathcal{T}|} \sum_{\mathcal{T}} (r_{ui} - q_i p_u)^2 + \left[\lambda_1 \sum_x \|p_u\|^2 + \lambda_2 \sum_x \|q_i\|^2 \right]$$

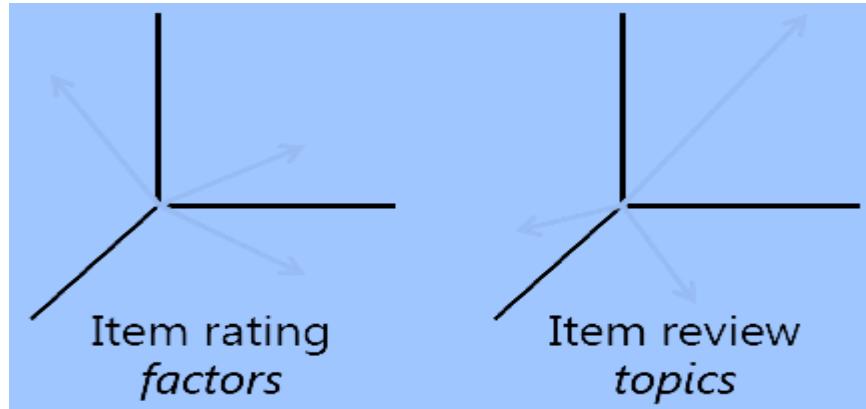
- Instead use review text as regularizer:

$$\min_{P,Q} \frac{1}{|\mathcal{T}|} \sum_{\mathcal{T}} (r_{ui} - q_i p_u)^2 - \lambda l(\mathcal{T} | \Theta, \phi, z)$$

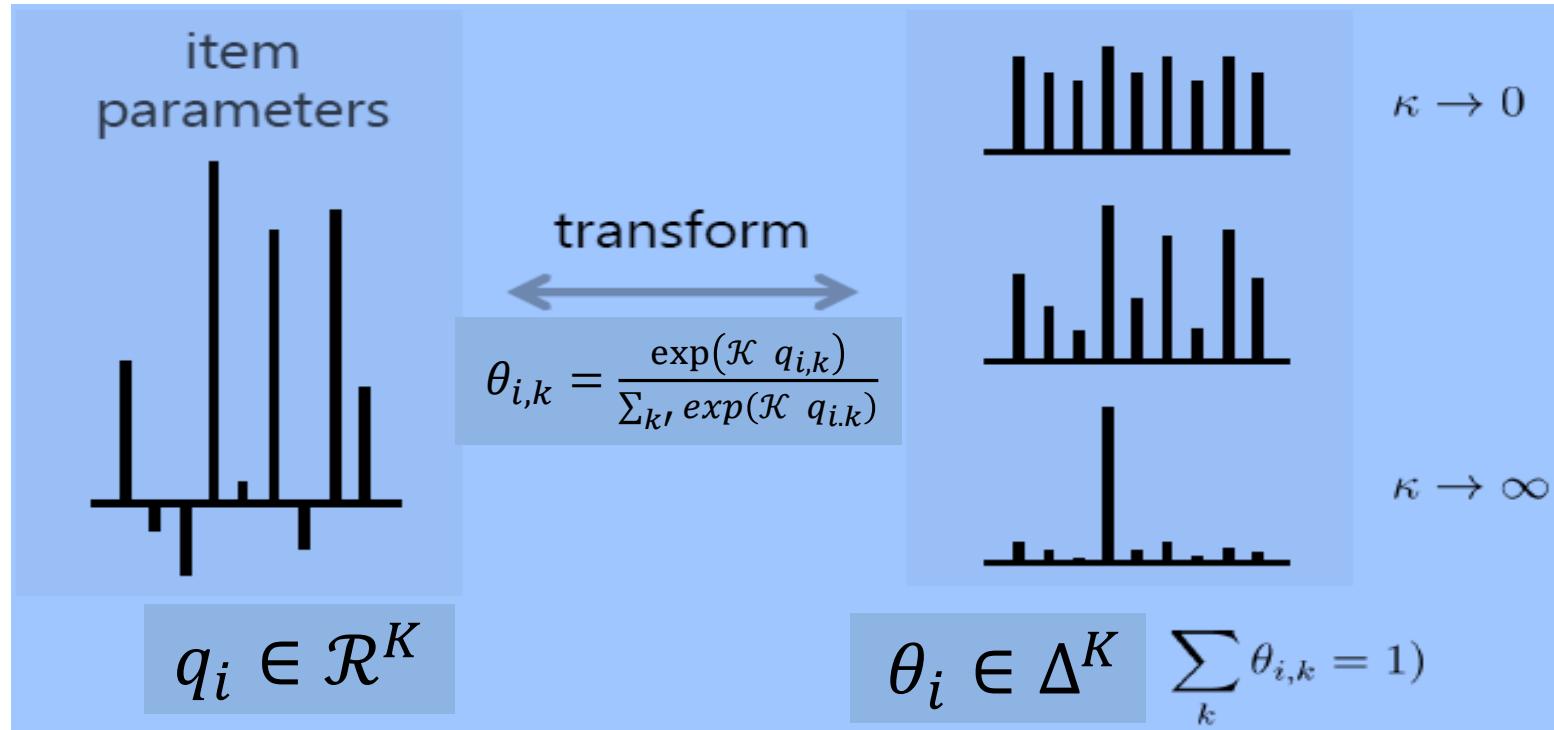
Rating error Corpus likelihood

HFT: Combining ratings and reviews

- Matrix factorization and LDA project users and items into low-dimensional space
- Align the two



HFT: Combining ratings and reviews



By linking rating and opinion models, find topics in reviews that inform about opinions

HFT: Model Fitting

Repeat Step 1 and Step 2 until convergence:

Step 1: Fit a rating model regularized by the topic
(solved via gradient ascent using L-BFGS)

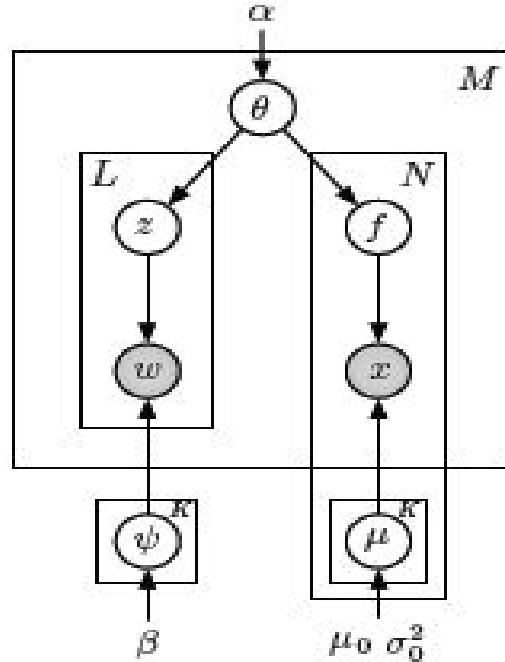
Step 2: Identify topics that “explain” the ratings.

Sample with probability
(solved via Gibbs sampling)

Ratings Meet Reviews, a Combined Approach to Recommend

- Use LDA to model reviews
 - Use mixture of Gaussians to model ratings
 - Combine ratings and reviews by sharing the same topic distribution
 - N users $\{u_i\}$
 - M items $\{v_j\}$
 - (u_i, v_j) defines the observed ratings $X = \{x_{ij}\}$,
 - Associated review r_{ij}
-
- Ratings meet reviews, a combined approach to recommend, Ling et al., Recsys 2014.

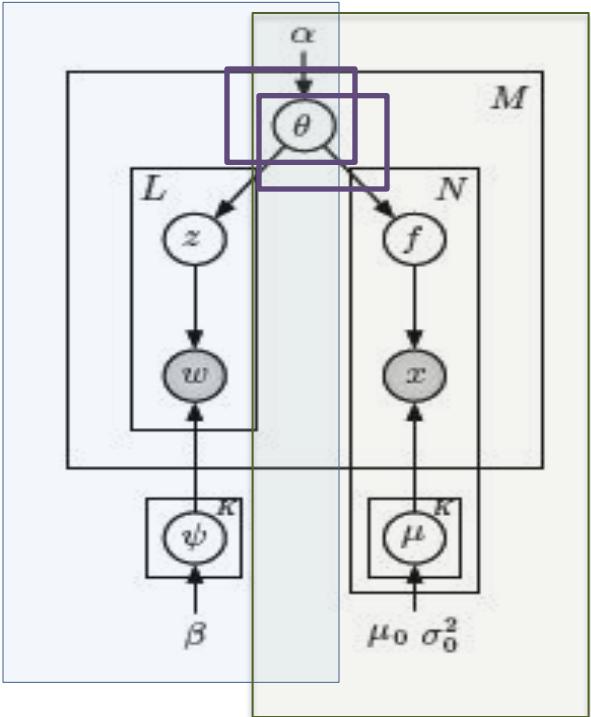
Graphical Model of RMR



$$P(\mathbf{w}, \mathbf{x} | \Theta; \alpha, \beta, \mu_0, \sigma_0^2, \sigma^2) \propto \prod_{j=1}^M P(\theta_j | \alpha) \prod_{i \in \mathcal{U}_j} \left(\prod_{l=1}^{m_j} \sum_{z=1}^n P(z | \theta_j) P(w_l | \psi_z) \right) \left(\sum_{f=1}^K P(f | \theta_j) P(x_{i,j} | \mu_{i,f}, \sigma^2) \right)$$

$$P(\mathbf{w}, \mathbf{x} | \Theta; \alpha, \beta, \mu_0, \sigma_0^2, \sigma^2) \propto \prod_{j=1}^M P(\theta_j | \alpha) \prod_{i \in \mathcal{U}_j} \left(\prod_{l=1}^{L_{i,j}} \sum_{z=1}^K P(z | \theta_j) P(w_l | \psi_z) \right) \left(\sum_{f=1}^K P(f | \theta_j) P(x_{i,j} | \mu_{i,f}, \sigma^2) \right)$$

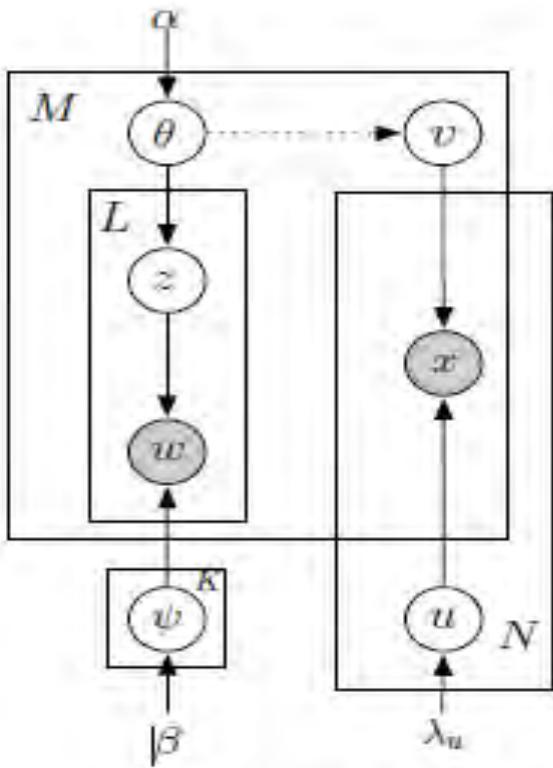
Use LDA to model the reviews



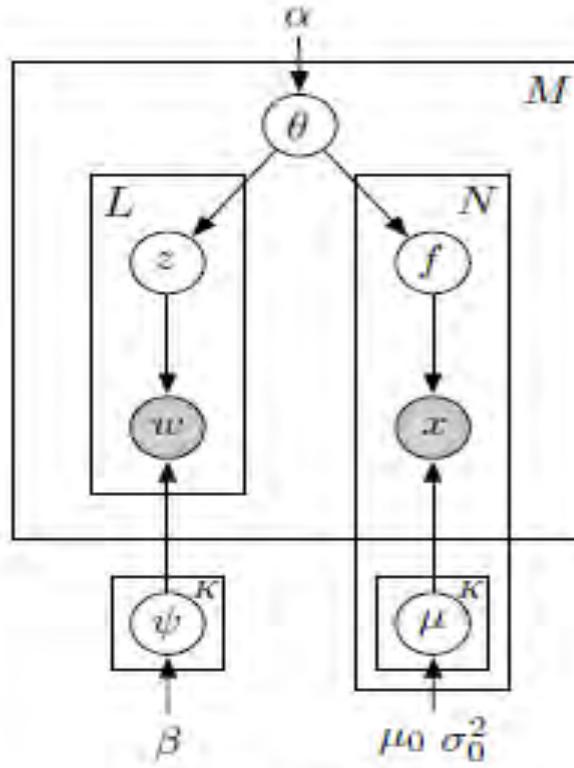
Use the same topic distribution to connect the rating part and the review part

Use mixture of Gaussians to model the ratings

Comparison of HFT and RMR



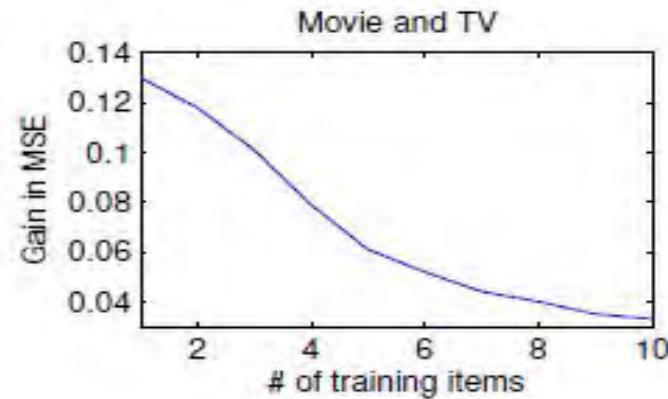
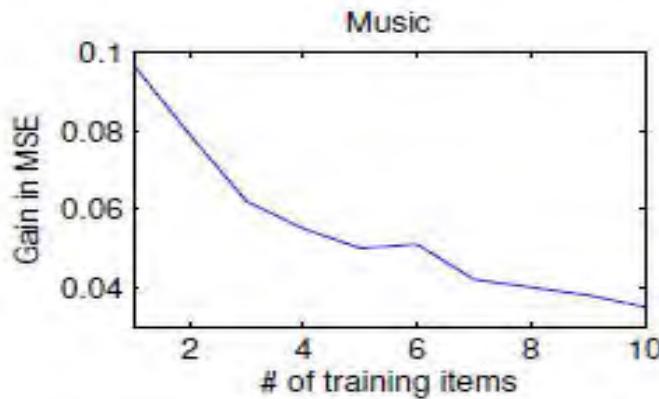
HFT



RMR

Results for Cold Start Setting

- The improvement of HFT and RMR over traditional MF is more significant for datasets that are sparse.



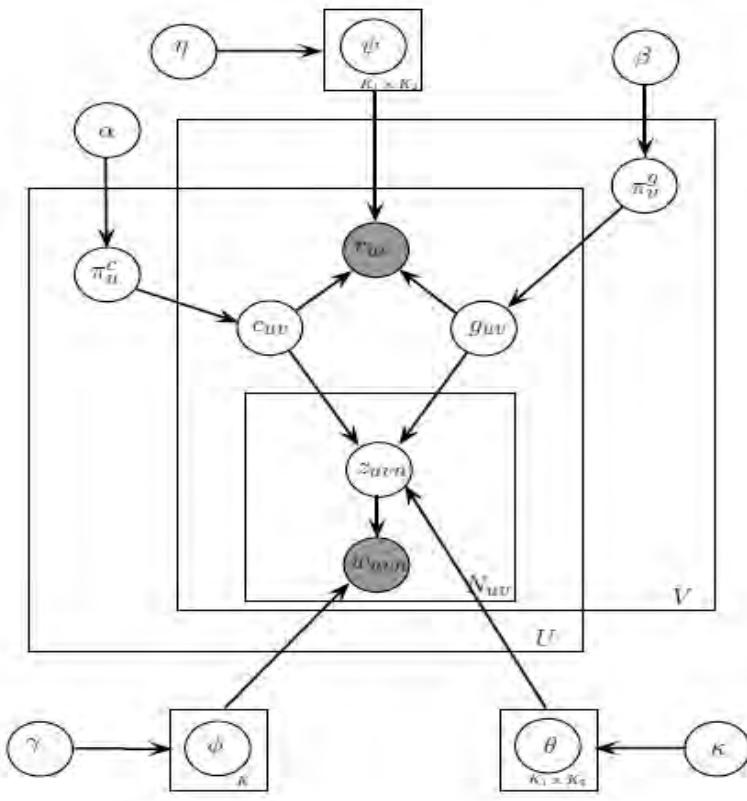
CMR

- Predict user's ratings by considering review texts as well as hidden user communities and item groups relationship.
- Uses Co-clustering to capture the relationship between users and items.
- Models as a mixed membership over community and group respectively.
- Each user or item can belong to multiple communities or groups with varying degrees.

Collaborative Filtering Incorporating Review Text and Co-clusters of Hidden User Communities and Item Groups, Yinqing Xu Wai Lam Tianyi Lin, CIKM 2014



CMR Graphical Model



U : user collection

V : item collection

r_{uv} : u's rating of item v

d_{uv} : u's review of item v

w_{uvn} : nth word of d_{uv}

z_{uvn} : topic of w_{uvn}

π_u^c : community membership of u

π_v^g : group membership of v

c_{uv} : user community of d_{uv} and r_{uv}

g_{uv} : item group of d_{uv} and r_{uv}

θ_{cg} : topic distribution of co-cluster (c,g)

ψ_{cg} : rating distribution of co-cluster (c,g)

ϕ_k : topic distribution of topic k

Collapsed Gibbs sampling algorithm to perform approximate inference.

TopicMF: Simultaneously Exploiting Ratings and Reviews for Recommendation

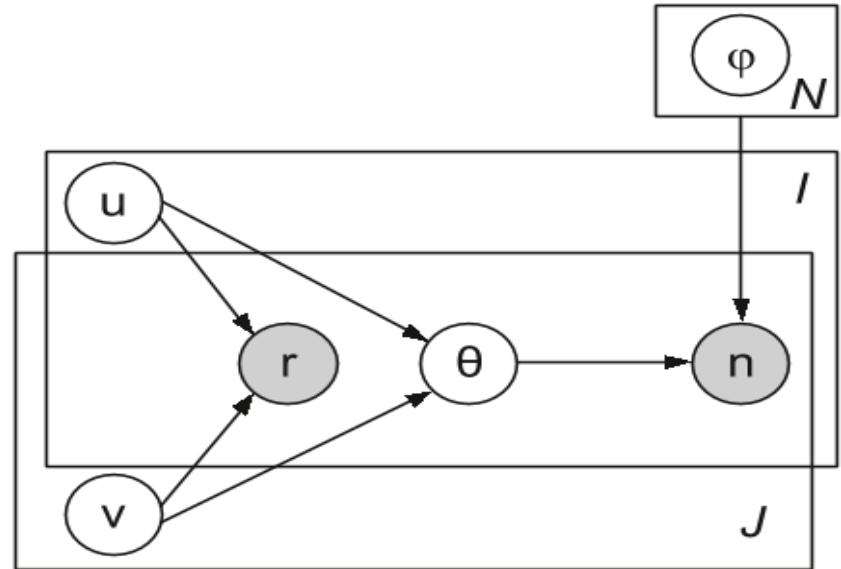
Yang Baot Hui Fang Jie Zhang, AAAI,2014

TOPIC-MF



TopicMF

- Matrix Factorization (MF) factorizes user-item rating matrix into latent user and item factors.
- Simultaneously topic modeling technique with Nonnegative Matrix Factorization (NMF) models the latent topics in review texts.

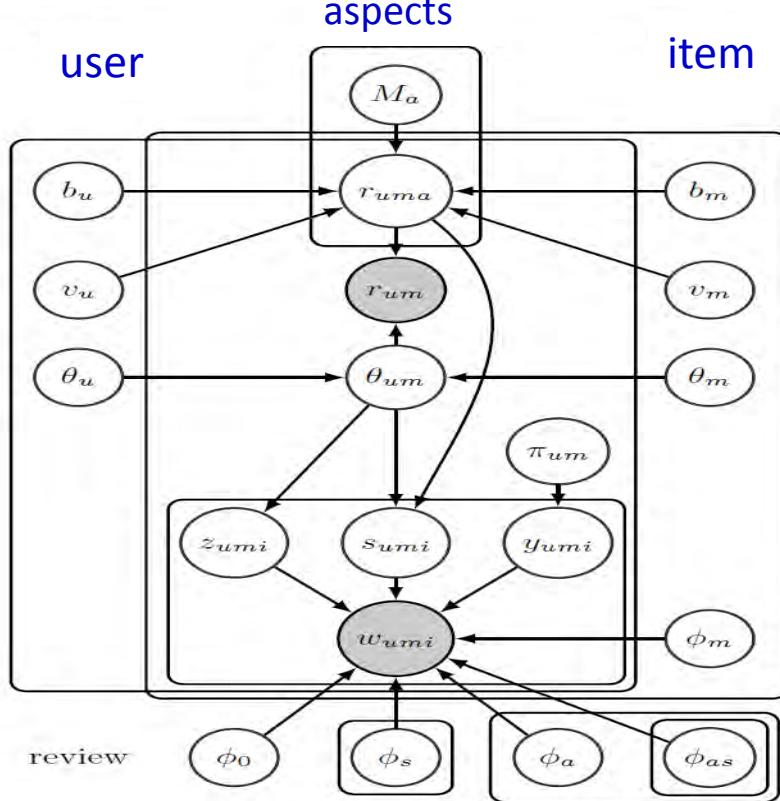


Align these two tasks by using a transform from item and user latent vectors to topic distribution parameters so as to combine latent factors in rating data with topics in user-review text.

TopicMF: Simultaneously Exploiting Ratings and Reviews for Recommendation, by Yang Bao, Hui Fang, Jie Zhang, AAAI, 2014

JMARS Model

Rating and Review Model



Aspect Modeling

(Interest of the user/ Property of the item)

$$\theta_u \sim \mathcal{N}(0, \sigma^2_{user-aspectI})$$

$$\theta_m \sim \mathcal{N}(0, \sigma^2_{item-aspectI})$$

$$\theta_{um} \sim \exp(\theta_u + \theta_m)$$

Rating Modeling

Aspect-specific rating

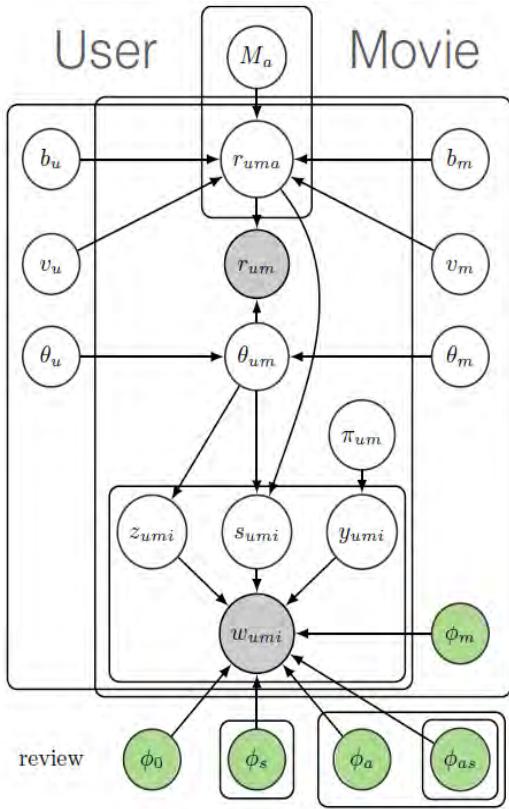
$$v_u^T M_\alpha v_m + b_u + b_m$$

Overall Rating

$$\hat{r}_{um} = \sum_a \theta_{uma} r_{uma}$$

M_α emphasizes the aspect specific properties.

JMARS Language Model



We assume that the review language model is given by a convex combination of five components.

- Background ϕ_0 :** words uniformly distributed in every review.
 - Sentiment ϕ_s :** not aspect specific content such as great, good, bad.
 - Item-specific ϕ_m :** any term that appears only in the item.
 - Aspect-specific ϕ_a :** words associated with specific aspects. Music, sound, singing
 - Aspect-specific sentiment words ϕ_{as} :** to express positive or negative sentiments.
- Each of the language models is a multinomial distribution.

Inference and Learning (hybrid sampling/optimization)

- The goal is to learn the hidden factor vectors, aspects, and sentiments of the textual content to accurately model user ratings and maximize the probability of generating the textual content.
- Use Gibbs-EM, which alternates between collapsed Gibbs sampling and gradient descent, to estimate parameters in the model.
- The Objective function consist of two parts :
 1. The prediction error on user ratings.
 2. The probability of observing the text conditioned on priors.

$$\mathcal{L} = \sum_{r_{um} \in \mathcal{R}} [\epsilon^{-2}(r_{um} - \hat{r}_{um})^2 - \log p(w_{um} | \Upsilon, \Omega)]$$

Qualitative Evaluation

To evaluate if the model is capable of interpreting the reviews correctly, the learned aspect ratings are examined.

Aspect rating, Sentiment and Movie-Specific Words

Aspect	Director	History	War	Life	Character
Rating	9.36	8.55	8.51	9.20	9.50
Prob	0.12	0.10	0.09	0.09	0.08

*... what an **excellent** piece of cinema ... the actors are great and directing **incredible** ... in **300**, Gerard Butler **dominates** the screen ... battle scenes are **incredible** ...*

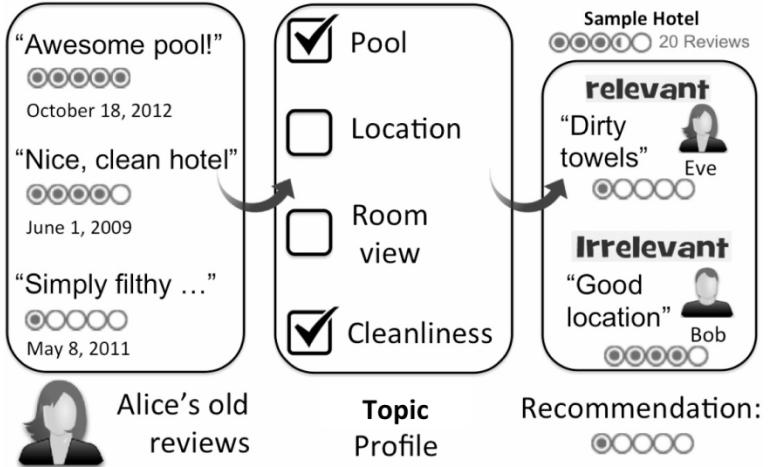
Table 4: The learnt aspect-specific ratings and latent sentiment identified by our model for a review.

Positive words

Movie-specific

TPCF

- Topic creation and manual verification and grouping.
- Faceted opinion extraction
- Based on the topics a user writes about, create an individual interest topic profile.
- Personalize the product rankings for each user, based on the reviews that are most relevant to her profile.



Recommendation Using Textual Opinions

Claudiu-Cristian Musat, Yizhong Liang, Boi Faltings, IJCAI 2013



Sentiment Utility Logistic Model

- SULM builds user and item profiles
 - Predicts overall rating for a review
 - Estimating sentiment utilities of each aspect.
- SULM identifies the most valuable aspects of future user experiences.
- Uses Opinion Parser for aspect extraction and aspect sentiment classification.
- SULM estimates the sentiment utility value for each of the aspects k using the matrix factorization approach.
- Tested on actual reviews across three real-life applications.

Aspect Based Recommendations: Recommending Items with the Most Valuable Aspects Based on User Reviews, by Konstantin Bauman, Bing Liu, Alexander Tuzhilin, KDD'17



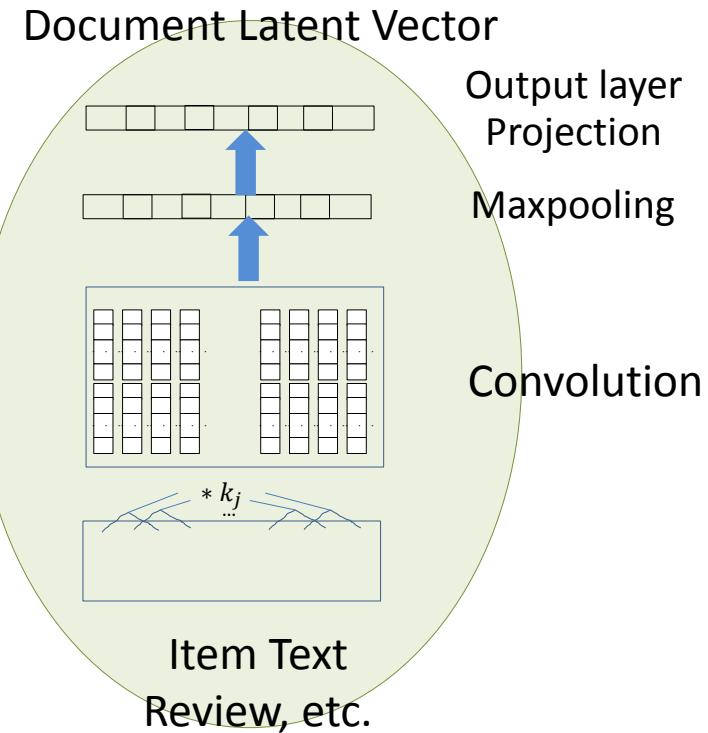
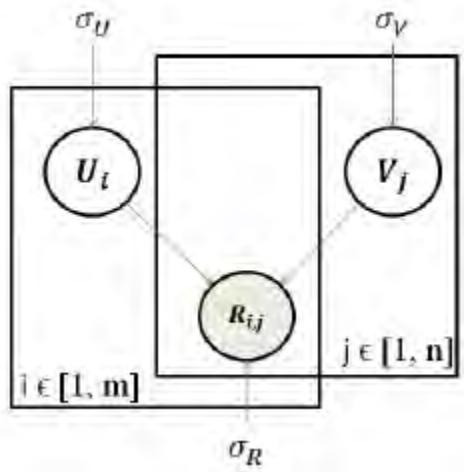
DEEP LEARNING BASED REVIEW AWARE RECOMMENDATION



CNN to find Document Latent Vector

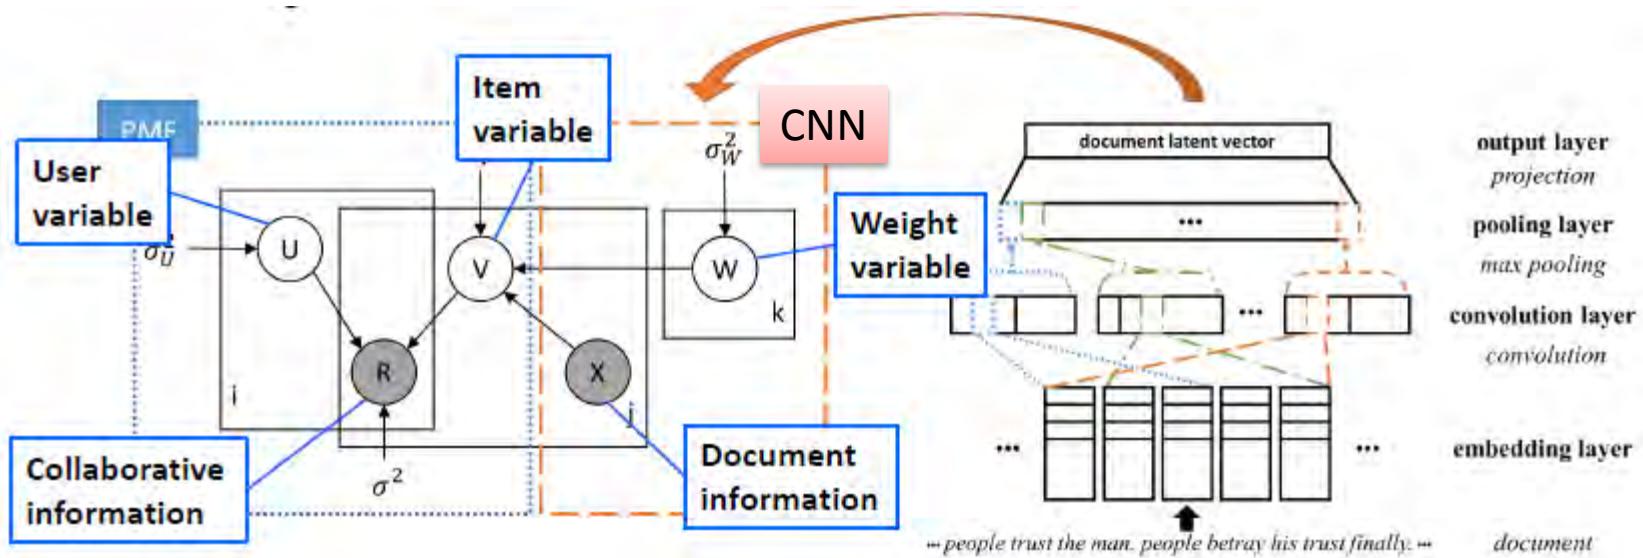
PMF

- CNN Architecture to generate document latent vector.
- Ratings approximated by probabilistic methods.



ConvMF

- Integrate CNN into PMF for the recommendation task.
- Item variable plays a role of the connection between PMF and CNN in order to exploit ratings and description documents.



Optimization

- Use Maximum a posteriori to solve U, V, W
- When U, V are temporarily fixed, to optimzie W, use backprop with given target variable v_j
- ConvMF+ used pretrained word embedding.

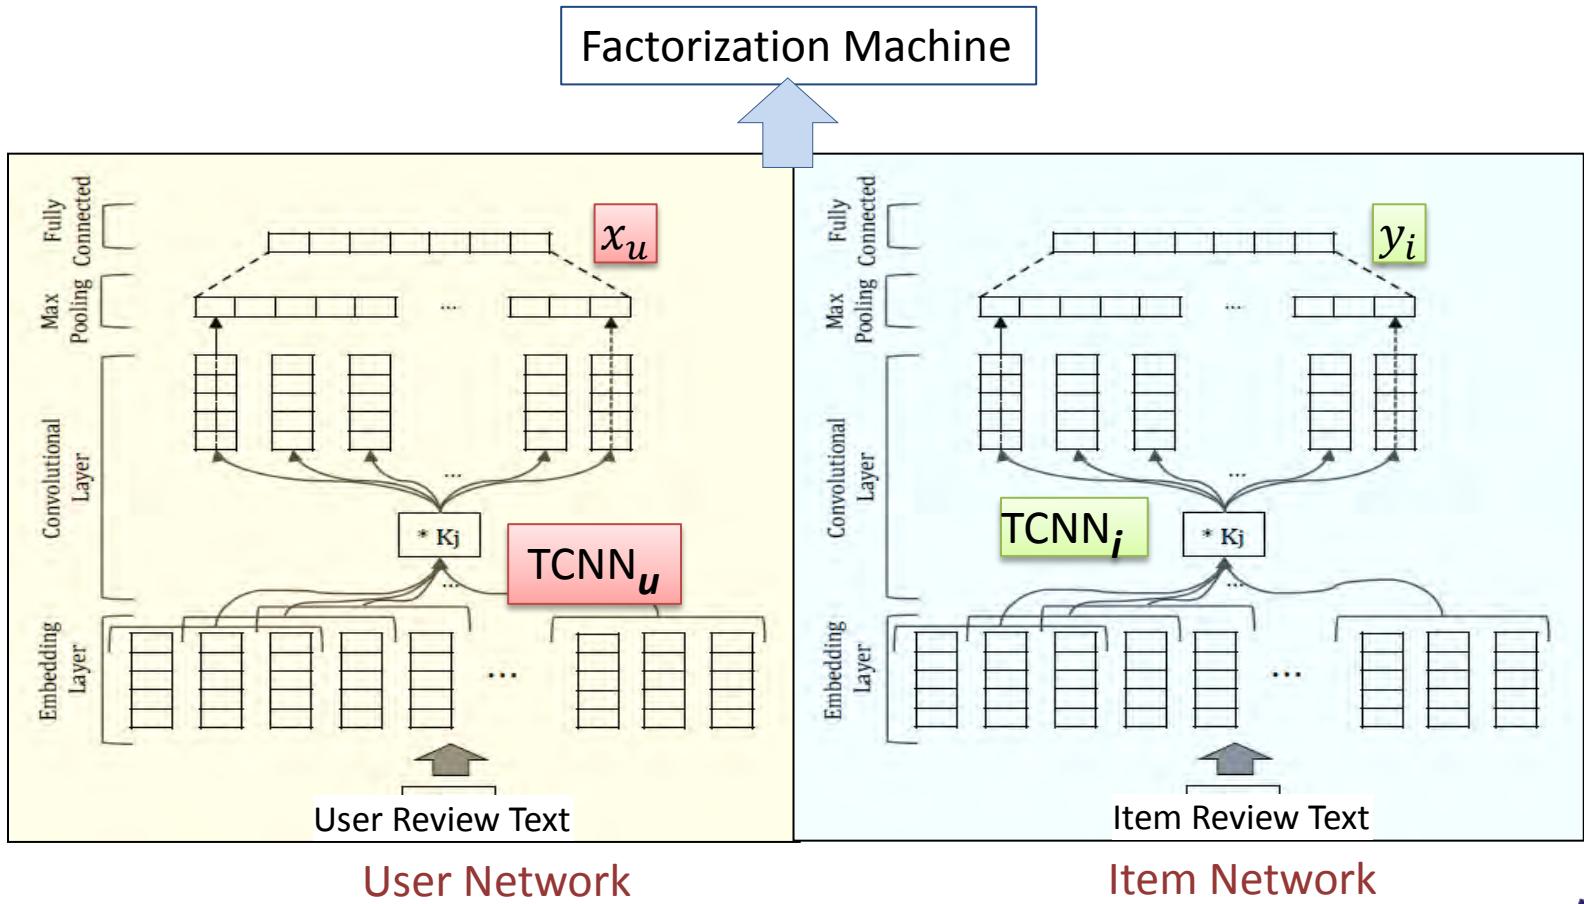
Deep Cooperative Neural Network

- (DeepCoNN) models users and items jointly.
- Two parallel convolutional neural networks to model user behaviors and item properties from review texts.
- Utilizes a word embedding technique to map the review texts into a lower-dimensional semantic space as well as keep the words sequence information. ↗
- The outputs of the two networks are concatenated.
- The factorization machine captures their interactions for rating prediction.

Joint deep modelling of Users and Items using Review for
Recommendation, Lei Zheng, Vahid Noroozi, Philip S. Yu, WSDM 2017



DeepCONN Architecture



User and Item Network

User Network

Lookup Layer:

- Reviews are represented as a matrix of word embeddings.
- All reviews written by user u are merged into a single document

$$v_{1:n}^u = \phi(d_1^u) \oplus \phi(d_2^u) \oplus \phi(d_3^u) \oplus \dots \oplus \phi(d_n^u)$$

Item Network

Lookup Layer:

- All reviews written on item i are merged into a single document

CNN Layers:

Each neuron j in the convolutional layer uses filter k_j on a window of words size t .

Max-pooling layer

$$o_j = \max\{z_1, z_2, \dots, z_{n-t+1}\}$$

Output from multiple filters:

$$O = \{o_1, o_2, \dots, o_{n1}\}$$

Fully connected layer: weight matrix W .

The output of the fully connected layer is considered as features of U .



Shared Layer

- A shared layer to map output of user and item network into the same feature space.

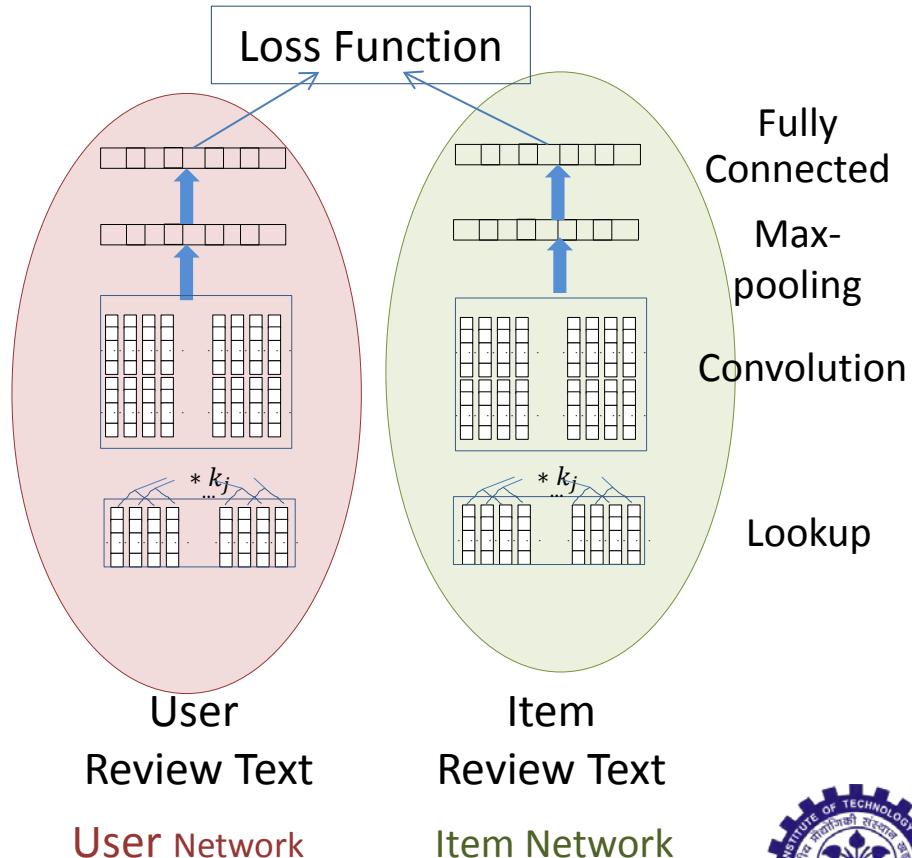
$$\hat{z} = (x_u, y_i)$$

- Factorization Machine(FM) computes the second order interactions between the elements of the input vector.

$$\hat{r}_{ui} = \hat{w}_0 + \sum_{i=1}^{\hat{z}} \hat{w}_i \hat{z}_i + \sum_{i=1}^{\hat{z}} \sum_{j=i+1}^{\hat{z}} \langle \hat{v}_i, \hat{v}_j \rangle \hat{z}_i \hat{z}_j$$

- L1 loss:

$$J = \sum_{(u,i,r_{ui}) \in D} |r_{ui} - \hat{r}_{ui}|$$



Experimental Results

Dataset	MF	PMF	LDA	CTR	HFT-10	HFT-50	CDL	DeepCoNN	Improvement of DeepCoNN (%)
Yelp	1.792	1.783	1.788	1.612	1.583	1.587	1.574	1.441	8.5%
Amazon	1.471	1.460	1.459	1.418	1.378	1.383	1.372	1.268	7.6%
Beer	0.612	0.527	0.306	0.305	0.303	0.302	0.299	0.273	8.7%
Average on all datasets	1.292	1.256	1.184	1.112	1.088	1.09	1.081	0.994	8.3%

TransNets: Learning to Transform for Recommendation
Rose Catherine William Cohen, RecSys 2017

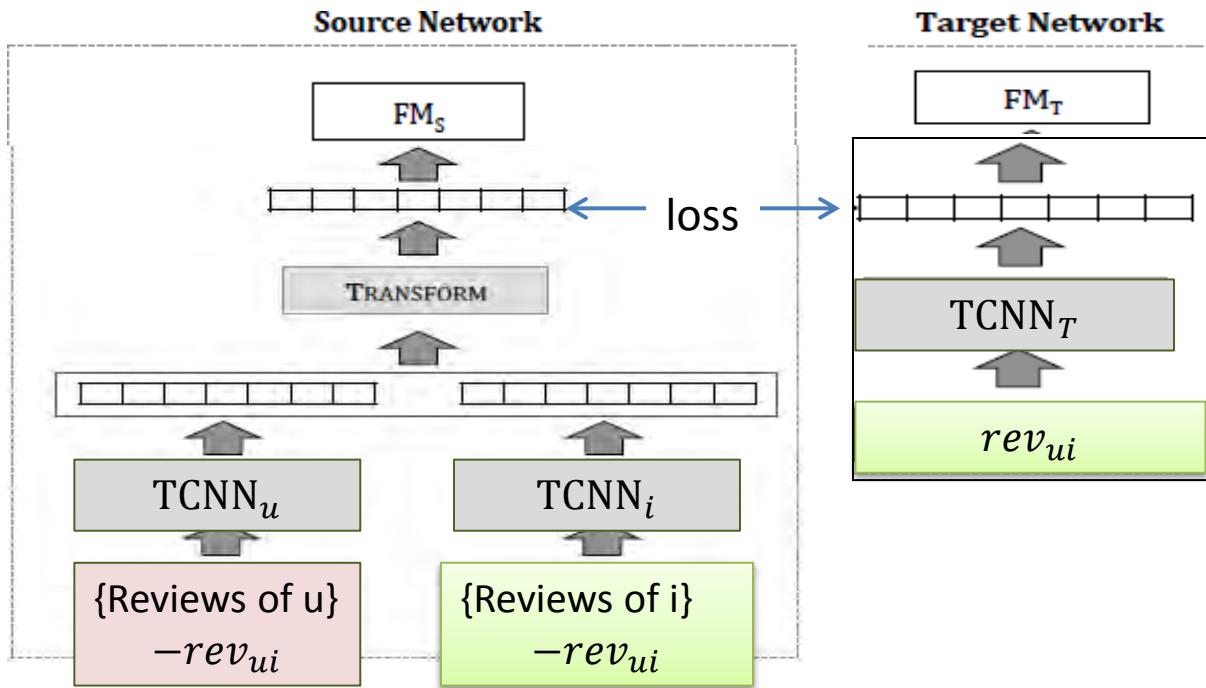
TRANSNETS



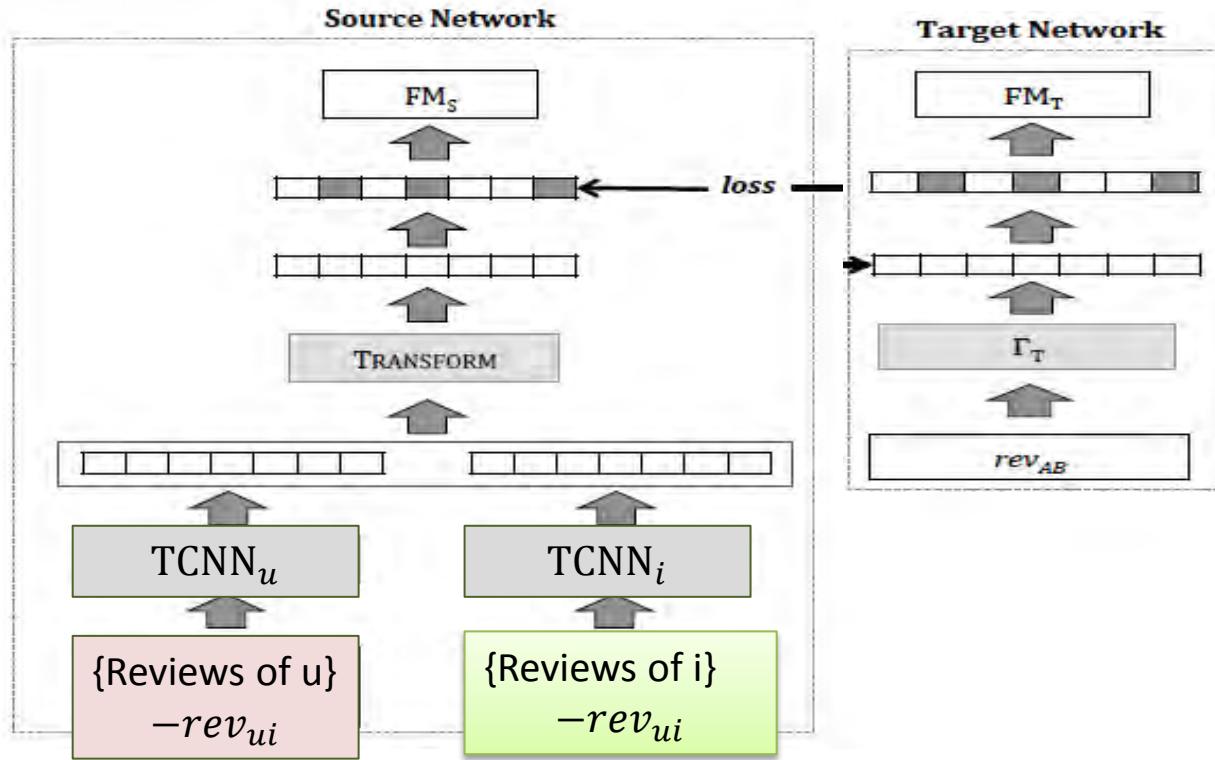
TransNets

- DeepCoNN performance is lower when the corresponding review is not in the training set.
- TransNets extends the DeepCoNN model by introducing an additional latent layer representing an approximation of the review corresponding to the target user-target item pair.
- Regularize this layer at training time to be similar to another latent representation of the target user's review of the target item.

TransNet



TransNets



Training TransNets

- During training, force the Source Network's representation z_L to be similar to the encoding of rev_{ui} produced by the Target Network.

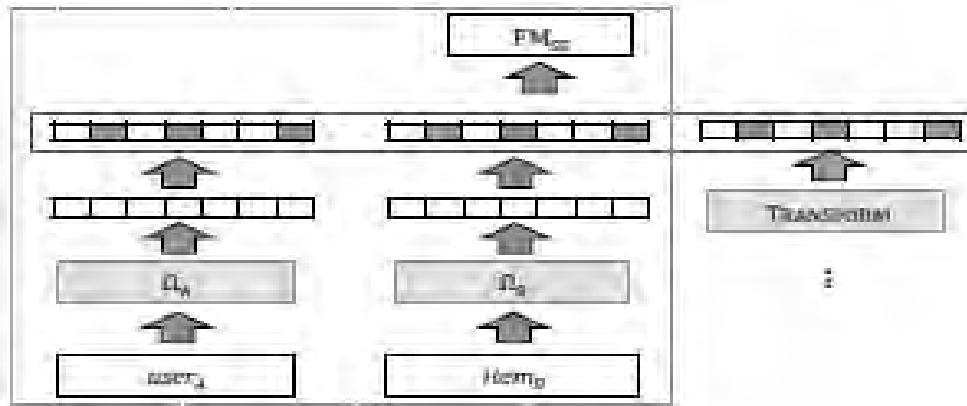
Repeat for each review:

- Step 1: Train Target Network on the actual review. Minimize L1 loss:
 $|r_{ui} - \hat{r}_T|$
- Step 2: Learn to Transform: minimize a L2 loss computed between z_L and the representation x_T of the actual review.
- Step 3: Train a predictor on the transformed input: Parameters of FM_S are updated to minimize a L1 loss : $|r_{ui} - \hat{r}_S|$

At test time, TransNet uses only the Source Network.

Extended TransNet

- TransNet does not use user, item identity.
- Ext-TN learn a latent representation of the users and items, similar to Matrix Factorization methods.
- Input to FM_S : latent user and item representation concatenated with the output of Transform layer.



Experimental Results

Performance comparison using MSE

Dataset	DeepCoNN + Test Reviews	MF	DeepCoNN	DeepCoNN-revAB	TransNet	TransNet-Ext
Yelp17	1.2106	1.8661	1.8984	1.7045	1.6387	1.5913
AZ-Elec	0.9791	1.8898	1.9704	2.0774	1.8380	1.7781
AZ-CSJ	0.7747	1.5212	1.5487	1.7044	1.4487	1.4780
AZ-Mov	0.9392	1.4324	1.3611	1.5276	1.3599	1.2691

D-Attn

 $\hat{r}_{u,i}$

dot product

FC layer

concatenation

FC layer

concatenation

max pool over sequence

convolution



local attention

 $W_1 \ W_2 \ \dots \ W_T$

L-Attn

max pool over sequence

convolution



global attention

 $W_1 \ W_2 \ \dots \ W_T$

G-Attn

User Network

max pool over sequence

convolution



local attention

 $W_1 \ W_2 \ \dots \ W_T$

L-Attn

max pool over sequence

convolution



global attention

 $W_1 \ W_2 \ \dots \ W_T$

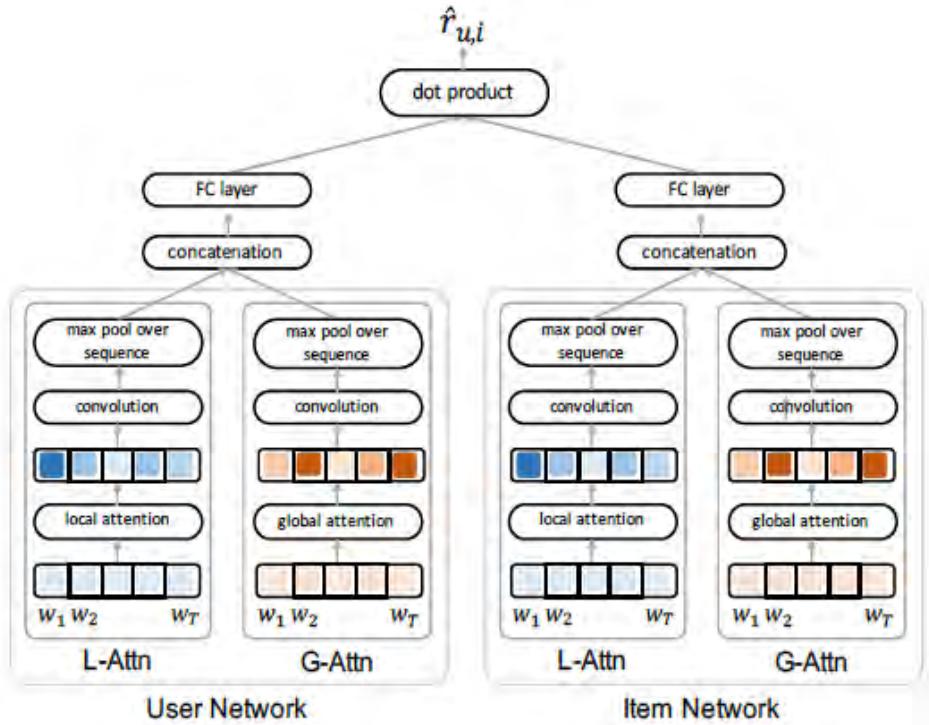
G-Attn

Item Network

Interpretable
Convolutional Neural
Networks with Dual
Local and Global
Attention for Review
Rating Prediction
Sungyong Seo, Jing
Huang, Hao Yang, Yan
Liu

D-Attn

- The local attention model (L-Attn) selects informative keywords from a local window.
- The global attention model (G-Attn) ignores irrelevant words from long review sequences.
- The outputs of L-Attn and G-Attn are concatenated and followed by a FC layer.
- Dot product of the vectors from both networks estimate ratings.



EXPLAINABILITY



Explainable Recommendation

- Why are these items recommended?

Your recently viewed items and featured recommendations

Inspired by your purchases



[Ms Draupadi Kuru: After the Pandavas](#)
Trisha Das
 18
Paperback



[Legends Of Halahala](#)
Appupen
 4
Paperback



[Maid in India: Stories of Inequality and...](#)
Tripti Lahiri
 4



[Ghost World](#)
Daniel Clowes
 4
Paperback

Reviews Explain Ratings

★★★★★ Use Different Charger

By RAKESHBS on 23 August 2017

Colour: Gold

Nice Phone ..

Pros:

- 1) Comfortable to hold in one hand
- 2) Good Camera Clarity
- 3) Finger print sensor works fast
- 4) Performance is good

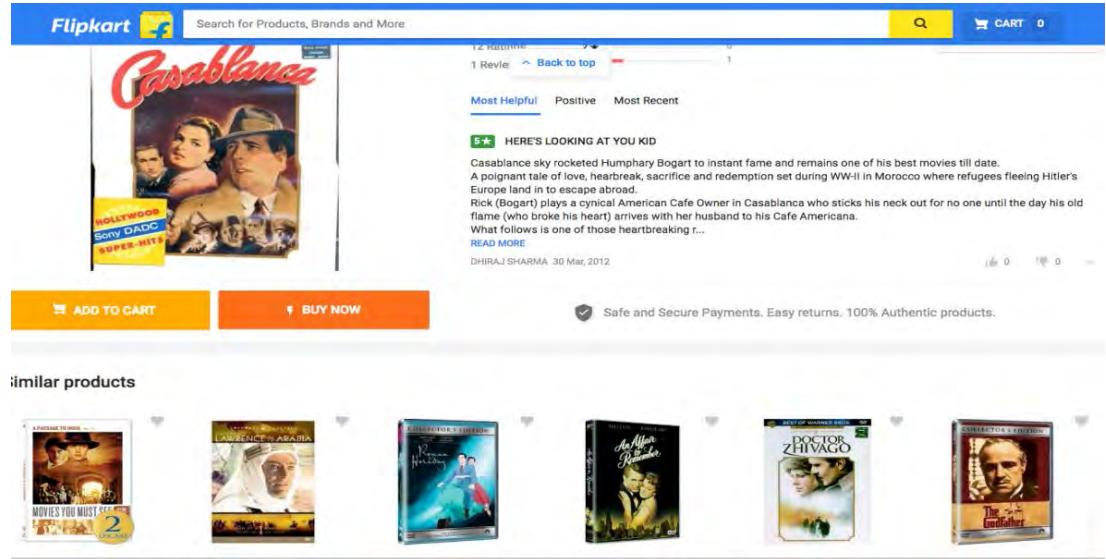
Cons:

- 1) it will get heated up when used with Charger shipped with Xiomi ..So I suggest to use other charger ...Like Nilkin 2A Charger or any other good 2A charger ...
- 2) Getting a good Tempered Glass will be costly and most of the time it is difficult to find one for 2.5D glass
- 3) Hybrid sim slot ...

- Reviews justify a user's rating:
 - by discussing the specific properties of items (aspects)
 - by revealing which aspects the user is most interested in.

Hidden factor/topics

- Interpretable textual labels for rating justification.
- Topics obtained explain variation in rating/review data



Hidden factors and hidden topics: understanding rating dimensions with review text,
Mcauley et al, RecSys 2013.

EFM

Explicit Factor Models for Explainable Recommendation based on
Phrase-level Sentiment Analysis

Yongfeng Zhang, Guokun Lai, Min Zhang, Yi Zhang, Yiqun
Liu, Shaoping Ma, SIGIR 2014

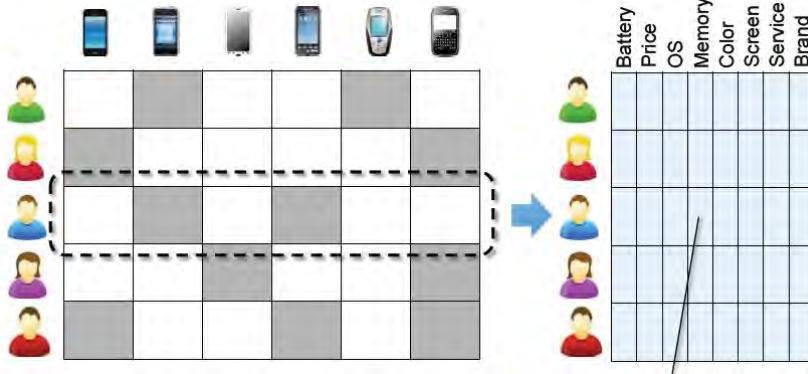


User-Feature Attention Matrix X

- Suppose feature F_j is mentioned by user u_i for t_{ij} times.

$$X_{ij} = \begin{cases} 0, & \text{if } u_i \text{ did not mention } F_j \\ 1 + (N - 1) \left(\frac{2}{1 + e^{-t_{ij}}} - 1 \right) & \text{otherwise} \end{cases}$$

- Rescale t_{ij} into the same range [1,N] as the rating matrix A by reformulating the sigmoid function.



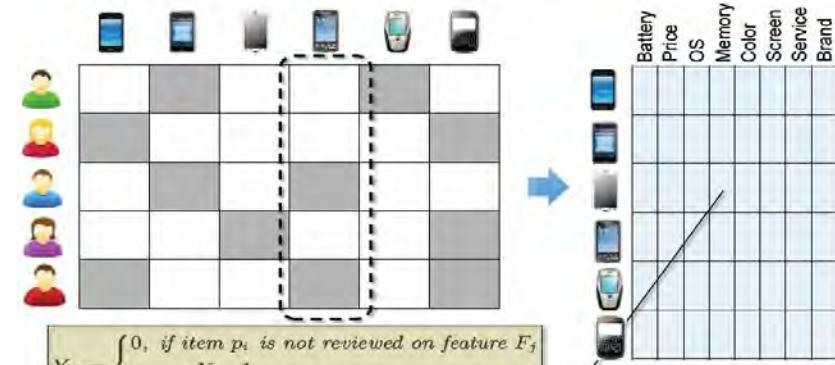
Item-Feature Quality Matrix Y

For each item p_i , extract the corresponding (F, S') pairs from all its reviews.

If feature F_j is mentioned for k times on item p_i , and the average of sentiment of feature F_j in those mentions are s_{ij} .

$$Y_{ij} = \begin{cases} 0, & \text{if item } p_i \text{ is not reviewed on } F_j \\ 1 + (N - 1) \left(\frac{2}{1 + e^{-k.s_{ij}}} - 1 \right) & \text{otherwise} \end{cases}$$

Rescaled to [1,N]



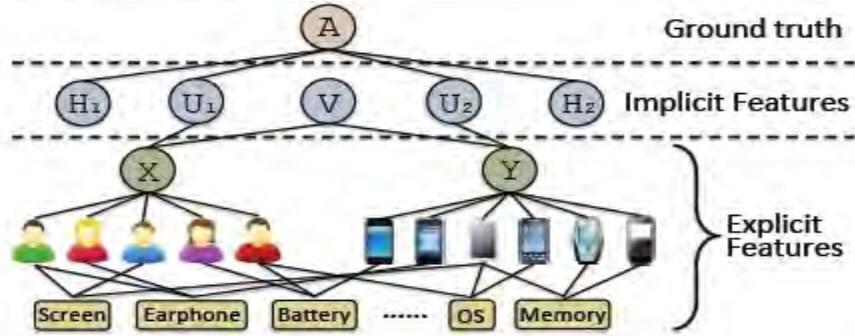
Factorization Model: Integrating Explicit and Implicit Features

- Build a factorization model over user-feature attention matrix X and item-feature quality matrix Y.

$$\underset{U_1, U_2, V}{\text{minimize}} \left\{ \lambda_x \|U_1 \cdot V - X\|_F^2 + \lambda_y \|U_2 \cdot V - Y\|_F^2 \right\}$$

- Introduce r' latent factors: H_1 and H_2 .
- We use $P = [U_1 H_1]$ and $Q = [U_2 H_2]$ to model the overall rating matrix A.
- U1: explicit factors
- H1: Hidden factors

$$\underset{P, Q}{\text{minimize}} \left\{ \|P \cdot Q - A\|_F^2 \right\}$$



$$\underset{U_1, U_2, V, H_1, H_2}{\text{minimize}} \left\{ \|P \cdot Q - A\|_F^2 + \lambda_x \|U_1 \cdot V - X\|_F^2 + \lambda_y \|U_2 \cdot V - Y\|_F^2 + \dots \right\}$$

Explanation Generation

- Template based explanation and an intuitional word cloud based explanation.
- Feature-level explanation for a recommended item.
- Provide disrecommendations by telling the user why the current browsing item is disrecommended.

You might be interested in <feature> on which this product performs well

You might be interested in <feature> on which this product performs poorly.

Trirank Approach to Aspect based Recommendation

Offline Training

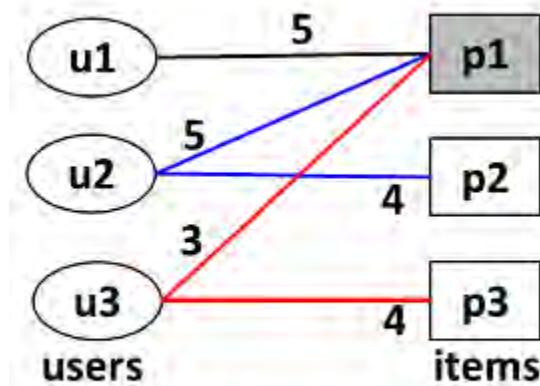
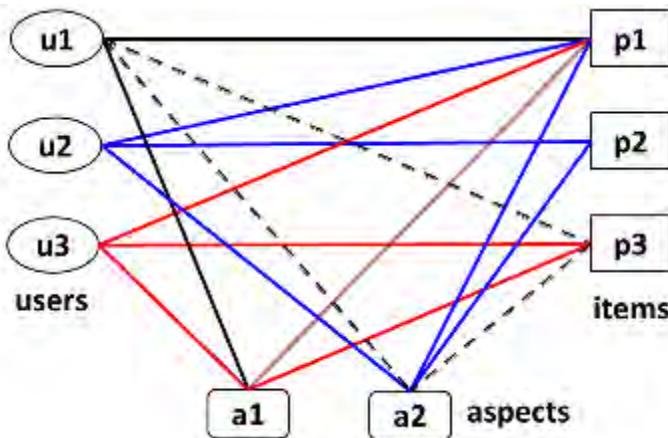
- Extract aspects from user reviews
- Build tripartite graph of User, Product, Aspect
- Graph Propagation: Label propagation from each vertex
- Machine learning for graph propagation

TriRank: Review-aware Explainable Recommendation by Modeling Aspects: Xiangnan He, Tao Chen, Min-Yen Kan, Xiao Chen

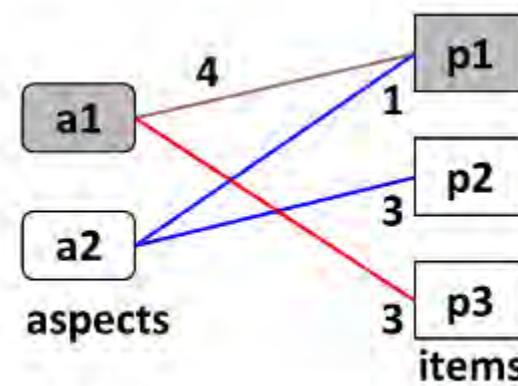


Inputs:

<u1, p1, a1>
<u2, p1, a2>
<u2, p2, a2>
<u3, p1, a1>
<u3, p3, a1>



(a) User-Item structure



(b) Item-Aspect structure

Trirank – Review Aware Recommendation

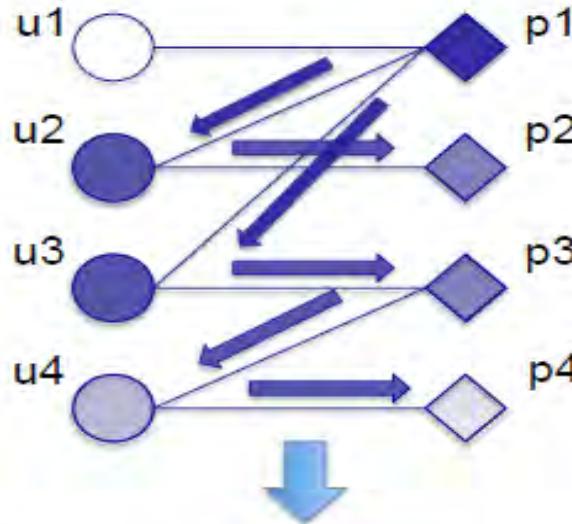
- Uses User-Item-Aspect ternary relation
 - Represented by heterogeneous tripartite graph
- Model reviews in the level of aspects
- Graph-based method for vertex ranking on tripartite graph accounting for
 - the structural smoothness (encoding collaborative filtering and aspect filtering effects)
Smoothness implies local consistency: that nearby vertices should not vary too much in their scores.
 - fitting constraints (encoding personalized preferences, prior beliefs).
- Offline training + online learning.
 - Provide instant personalization without retraining.

Basic Idea: Graph Propagation

Inputs:

<u1, p1, l>
<u2, p1, l>
<u2, p2, l>
<u3, p1, l>
<u3, p3, l>
<u4, p3, l>
<u4, p4, l>

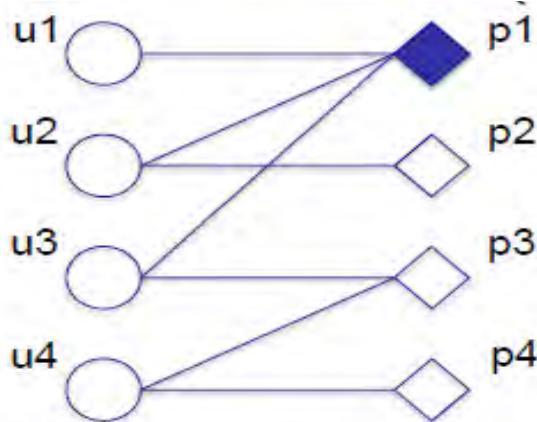
Target user: u1



Item ranking: $p_2 \approx p_3 > p_4$
User ranking: $u_2 \approx u_3 > u_4$

Label Propagation from the target user's historical item nodes captures
Collaborative Filtering

Machine Learning for Graph Propagation



- Input:
 - Graph Structure (Matrix Y)
 - Initial labels to propagate (vectors p^0)
- Output:
 - Scores of each vertex (vectors u, p)

Smoothness Regularizer: Nearby vertices should not vary too much

$$\sum_{i \in U} \sum_{j \in P} y_{ij} \left(\frac{u_i}{\sqrt{d_i}} - \frac{p_j}{\sqrt{d_j}} \right)^2$$

Fitting Constraints (initial labels):
Ranking scores should adhere to the initial labels.

$$\sum_{j \in P} (p_j - p_j^0)^2$$

Optimization (Gradient Descent):

$$p = S_Y u + p^0$$

$$u = S_Y^T p, \quad \text{where } S_Y = \left[\frac{y_{ui}}{\sqrt{d_u d_i}} \right]$$

[He et al, SIGIR 2014]

Connection to CF models

- Recap: ranking loss function (for a target user):

$$\boxed{\sum_{j \in P} (p_j - p_j^0)^2} + \lambda \sum_{i \in U} \sum_{j \in P} y_{ij} \left(\frac{u_i}{\sqrt{d_i}} - \frac{p_j}{\sqrt{d_j}} \right)^2$$

Prediction loss Regularizations

- Traditional machine learning-based CF models:

- #### I. Prediction model:

E.g., matrix factorization: $\hat{y}_{ui} = \langle v_u, v_i \rangle$

- ## 2. Loss function:

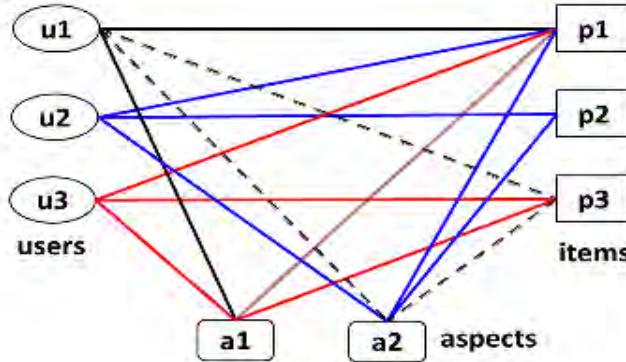
$$\sum_{u \in U} \sum_{i \in I} (y_{ui} - \hat{y}_{ui})^2$$

Prediction loss on all items (include imputations).
(important for top-K recommendation)

Trirank Approach

- Graph propagation in the tripartite graph:

Inputs:
 $\langle u_1, p_1, a_1 \rangle$
 $\langle u_2, p_1, a_2 \rangle$
 $\langle u_2, p_2, a_2 \rangle$
 $\langle u_3, p_1, a_1 \rangle$
 $\langle u_3, p_3, a_1 \rangle$



$$\mathbf{u} = \mathbf{u}^0 + \lambda_1 UP \cdot \mathbf{p} + \lambda_2 UA \cdot \mathbf{a}$$

$$\mathbf{p} = \mathbf{p}^0 + \lambda_3 PU \cdot \mathbf{u} + \lambda_4 PA \cdot \mathbf{a}$$

$$\mathbf{a} = \mathbf{a}^0 + \lambda_5 AP \cdot \mathbf{p} + \lambda_6 AU \cdot \mathbf{u}$$

Initial labels should encode:

- Target user's preference on aspects/ items/ users:
 - a^0 : reviewed aspects.
 - p^0 : ratings on items.
 - u^0 : similarity with other users (friendship).

Online Learning

Offline Training (for all users):

1. Extract aspects from user reviews
2. Build the tripartite graph model (edge weights)
3. Label propagation from each vertex and save the scores.
 - *i.e.* store a $|V| \times |V|$ matrix $f(v_i, v_j)$.

Online Recommendation (for target user u_i):

1. Build user profile (*i.e.*, L_u vertices to propagate from).
2. Average the scores of the L_u vertices:

$$y_j = \frac{1}{|L_u|} \sum_{v_u \in L_u} f(v_u, v_j)$$

Explainability

- Collaborative filtering + Aspect filtering

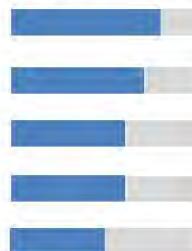
(Similar users also choose the item)

(Reviewed aspects match with the item)

- Item Ranking
- Aspect Ranking
- User Ranking

Chick-Fil-A is recommended for you based on your preference on its aspects.

Speciality ↓



Your Preference



Dislike the recommendation? Change your preference [here!](#)

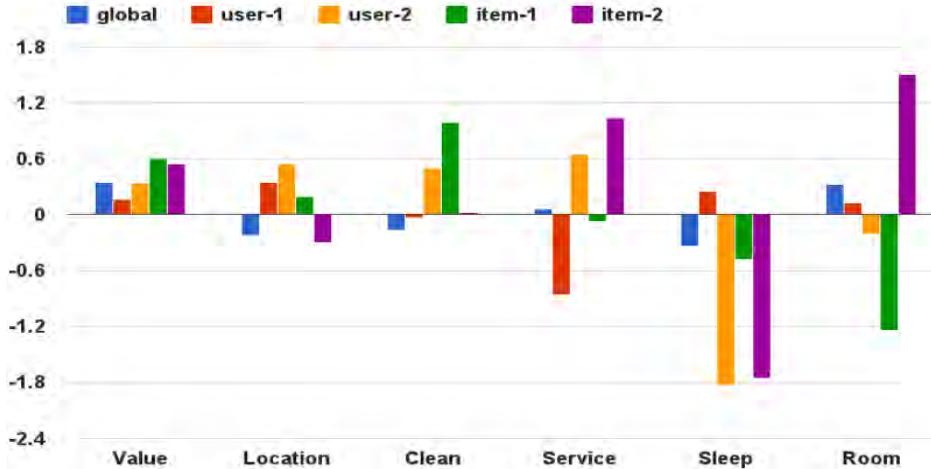


FLAME

- FLAME learns users' personalized preferences on different aspects from their past reviews, and predicts users' aspect ratings on new items by collective intelligence.
- Propose a new model combining aspect-based opinion mining and collaborative filtering to collectively learn users' preferences on different aspects.
- Preference diversity: For example, the food of the same restaurant might be delicious for some users but terrible for others.
- To address the problem of Personalized Latent Aspect Rating Analysis, we propose a unified probabilistic model called Factorized Latent Aspect ModEl (FLAME), which combines the advantages of both collaborative filtering and aspect-based opinion mining

FLAME: Probabilistic Model Combining Aspect Based Opinion Mining and Collaborative Filtering, Yao Wu and Martin Ester, WSDM'15

Aspect Weights, Aspect Values



- user-1 likes to comment on the Location and Sleep
- user-2 cares more about the Service, Clean and Location.

Produce more persuasive recommendation explanations by the predicted aspect ratings and some selected reviews written by similar users.



(a) Location



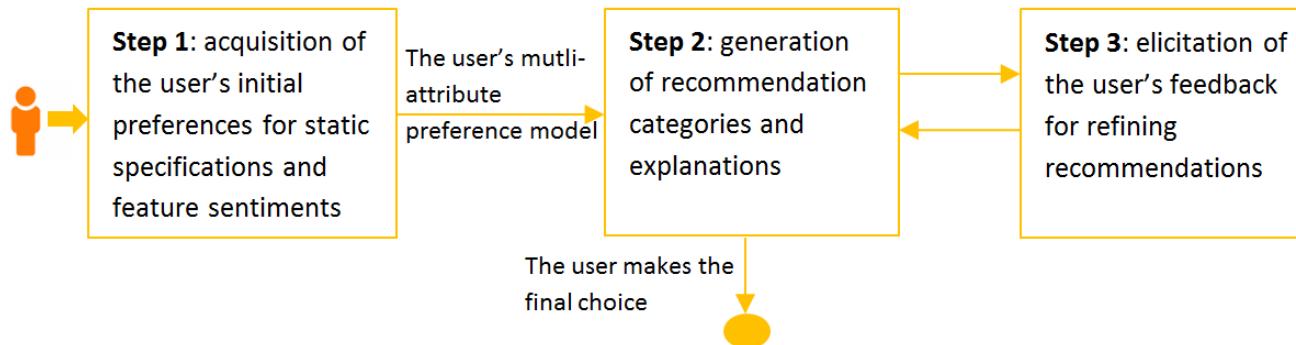
(b) Location 2-star



(c) Location 5-star

Senti-ORG: Explanation Interface

- An explanation interface that emphasizes explaining the tradeoff properties within a set of recommendations in terms of both their static specifications and feature sentiments extracted from product reviews.
- Assist users in more effectively exploring and understanding product space.



Explaining Recommendations Based on Feature Sentiments in Product Reviews
Li Chen, F Wang, IUI 2017



Samsung Galaxy Camera (Wi-Fi)



Avg. Rating: ★★★★☆ 4.2 (22 reviews)

Price:	\$385.0
Screen Size:	4.8 inches
Effective pixels:	16.0 megapixels
Optical zoom:	20.9 x
Weight:	0.88 pounds

[More Details](#) [Add to Saved List](#)

The top candidate

Related Cameras

They have **better values at effective pixels, weight, price, but worse value at screen size**

Sony Cyber-shot DSC-HX20V



★★★★☆ 4.4
(189 reviews)

Price: \$299.0
Screen Size:
Effective pixels:
Optical zoom:
Weight:

They have **better values at optical zoom, price, but worse value at weight**

Panasonic Lumix DMC-FZ60

Price: \$349.0

Tradeoff-oriented category explanation based on static specifications (e.g., “**They have better values at effective pixels, weight, price, but worse value at screen size**”)

Sony Cyber-shot DSC-TX200V



Price: \$299.0
Screen Size: 3.3 inches
Effective pixels: 18.0 megapixels

(120 reviews)

Nikon Coolpix P510



Price: \$303.5
Screen Size: 3.0 inches
Effective pixels: 16.0 megapixels

Future Trends

- Explanation Interfaces
 - User control on recommendations.
 - Use with other knowledge sources: context, ontology,
...
- Use of reviews beyond recommendations.
- Addressing Complex and Subjective Product-Related Queries, (McAuley WWW'16)

Thank You

